

CS 351 Computer Graphics, Fall 2014

Dr. Bruce A. Maxwell
Department of Computer Science
Colby College

Course Description

Computer graphics deals with the manipulation and creation of digital imagery. We cover drawing algorithms for two-dimensional graphics primitives, 2D and three-dimensional matrix transformations, projective geometry, 2D and 3D model representations, clipping, hidden surface removal, rendering, hierarchical modeling, shading and lighting models, shadow generation, special effects, fractals and chaotic systems, and animation techniques. Labs will focus on the implementation of a 3D hierarchical modeling system that incorporates realistic lighting models and fast hidden surface removal.

Prerequisites: CS 251 or permission of instructor. Linear algebra recommended.

Desired Course Outcomes

- A. Students understand and can implement the fundamental concepts of creating and manipulating images.
- B. Students understand and can implement the fundamental concepts of rendering, including matrix transformations, shading, and hidden surface removal.
- C. Students understand and can implement the fundamental concepts of modeling objects and scenes and hierarchical modeling systems.
- D. Students work in a group to design and develop 3D modeling and rendering software.
- E. Students present methods, algorithms, results, and designs in an organized and competently written manner.
- F. Students write, organize and manage a large software project.

This material is copyrighted by the author. Individuals are free to use this material for their own educational, non-commercial purposes. Distribution or copying of this material for commercial or for-profit purposes without the prior written consent of the copyright owner is a violation of copyright and subject to fines and penalties.

1 Graphics

If you could make a picture, what would it be?

1.1 Graphics is...

The field of computer graphics impacts many disciplines and activities. In some cases its use is specialized to a particular application, such as scientific visualization. In other cases, its impact is broad and ubiquitous, such as windowing systems on a computer.

- Vector graphics: **rendering** images using points and lines.

(Rendering: to create the actual image, to make it visible to the user. "cause to be or become; make.")

Some types of images are built from lines and shapes. Common examples include logos, fonts, and wire-frame models. An image like a logo or a letter should look the same regardless of its size. The important characteristics of these images are not their pixels, but the geometric relationship of the image elements, which can be described as a set of points, lines, vectors, arcs, polygons, or other shapes. If we represent these types of images as geometric entities, converting them to pixels only as the image is being rendered to a device, then we are free to rotate, scale, translate, and otherwise transform them arbitrarily, manipulating only the points and vectors being used to represent them. If we had to represent a letter as a set of pixels, then scaling it to a larger size would result in blocky artifacts (the **jaggies**).

- 2-D shapes and shape manipulation: presentation software

When we are creating documents we often want to manipulate the elements of the document in arbitrary ways. Again, we are best served by thinking about the elements of the document as mathematical objects described by points and vectors, which are amenable to mathematical transformations. Selecting, moving, scaling, rotating, and transforming these elements are all enabled by the fundamental concepts of CG.

- Drawing and photo manipulation programs: from MacPaint to Photoshop

Drawing and photo manipulation programs often blend vector graphics with pixel-based manipulation. Manipulating pre-existing photos is a branch of computer graphics focused on both creative manipulation and post-processing for photo enhancement. Much of the field of photo manipulation falls in the category of image processing, or signal processing more generally, where we think of the image as a 2-D signal or function to be filtered and modified.

- Graphical User Interfaces [GUI] and computer windowing interfaces

Graphical user interfaces are the primary method we have of interacting with most computers (e.g. windowing systems, widgets, buttons and sliders, point and click, click and drag). Rendering the windows in the proper order and ensuring the user sees the correct information are core computer graphics concepts. Clipping—identifying when one graphic item occludes another—is a key concept underlying windowing systems, and fast clipping algorithms are critical to making GUIs run fast enough to be practical. Likewise, all of the various effects from bouncing applications to sliding buttons to flowing windows are the result of CG algorithms.

- **Rendering** pictures of 3-D scenes from models

Taking a 3-D model and creating a 2-D image of it from a selected viewpoint is the canonical computer graphics problem. It encompasses model representation and creation, the definition of a view, the definition of the viewing geometry, the projection of a 3-D world onto the 2-D viewing geometry, and the rendering of geometric objects on to a raster representation. In each stage, the efficiency of the representation and the speed of the process are critical to making CG practical.

- **Animation** of models

Making things appear to move realistically is hard to do. While rigid objects can be usually be controlled through Newtonian mechanics, once things have multiple jointed connections, become floppy or springy, can fall apart, can interact with other objects, or deform under a variety of forces, the equations and interactions can become extremely complex. In addition, when making things move we often want them to be in certain places at certain times, such as when animating a character in a movie. How do we make the motion natural, while still achieving a specific outcome? There are a large number of tools available not only for creating motions from scratch, but also for capturing the motion of real things and then mapping that motion onto virtual objects. After rendering, animation is the second fundamental task in CG.

- **Modeling** of physical or virtual phenomena

In addition to animating objects, modeling real-world or other-world phenomena—static or dynamic—is the third fundamental task of CG. Examples include creating realistic wear on stone monuments, the accurate modeling of reflections on shiny surfaces or sub-surface scattering in skin or semi-transparent liquids, and realistic smoke, fire, and water. There are many visual phenomena that are difficult to model and require complicated algorithms or equations to simulate properly.

- Computer-Aided Design [CAD] or Computer-Aided Modeling [CAM]

Almost all manufactured devices available today were initially designed and modeled using a computer graphics program. CAD design is any process where a computer is the primary design tool, and computer-aided modeling enables testing of device or object functionality prior to physically creating the object. With the advent of 3D printers, it is now possible to create fast prototypes, or final versions, of almost any design in many different types of materials.

- Training and simulation, such as flight simulators

Training people for a complex task can be expensive, if not outright dangerous. Flight simulators were one of the first commercial applications of computer graphics because training people in real airplanes, and putting those airplanes through dangerous maneuvers is an expensive proposition (time, money, and lives). One common application of CG is to provide useful, often realistic, interactive training systems without requiring real world resources and exposing the trainees to dangerous situations.

- Virtual reality and immersive 3D environments

Useful for training, research, and entertainment, virtual reality [VR] and immersive 3D environments enable people to experience a more complete version of a 3D world. there are many different types of VR systems with a range of capabilities. But the key difference between a VR system and a passive system is that the system must respond to the motions of the viewer to maintain the illusion that the viewer is inside the virtual world and their motions and actions have the correct effects on that world.

- Visualization of data, such as medical imaging or web networks

How to visualize data in ways that enable people to see patterns is an important sub-field of CG. There are a variety of alternative projection schemes that map data into different types of spaces for visualization. A simple example is visualizing data on a sphere. We can unwrap the sphere to create a 2-D map, or we can try to create an interactive visualization of the data on a sphere that can be moved and rotated by the user so that it appears to retain its 3D shape—even though it is being viewed on a 2D image.

- Games, both 2D and 3D

Computer games are now the largest segment of the visual entertainment industry. Almost all visual aspects of games fall into the realm of CG, including the realistic modeling and AI systems necessary to create challenging opponents that move and act appropriately. (Note the blurry lines between CG and other domains when modeling things in the world starts to require a wide variety of knowledge and algorithms from other disciplines.)

1.2 Representing images and colors

At its core, graphics is about creating images from some kind of input data. A video is simply a coherent sequence of images. Images consist of a combination of spectral and spatial information. In vector graphics, the information is encapsulated as line segments or other crisp geometric entities. In raster graphics, the image is represented by a set of pixels. Normally, the pixels are laid out on a grid, such as a regular rectangular grid, and each pixel is represented by a 3-tuple (x, y, \vec{C}) , where \vec{C} is the spectral information. In most consumer devices and sensors, the spectral information consists of a 3-tuple (R, G, B). However, \vec{C} may be a scalar value for greyscale devices or hundreds of values for hyper-spectral imaging devices.

What is color?

One important concept to understand as a CG developer or artist is that color is a perception. The perception of color is caused by electromagnetic [EM] radiation hitting sensing cells in our eyes. While different EM spectra will cause different responses in our sensors, it is not the case that different EM spectra will always cause different perceptions of colors. In fact, it is possible to adjust a person's color perception of a particular EM spectrum by modifying the EM spectra of surrounding areas. This has been shown in many psychophysical studies, and CG can take advantage of this fact in order to create particular experiences for the user.

For example, if we want to make the user think there is a light source in the scene, then we have to control the surrounding context to make the light source appear brighter than what is around it. Given that we are limited to a fixed gamut of colors, and the screen is limited to a maximum amount of energy at each pixel, the only way to make something appear to be emitting light is to make the rest of the scene darker. The amazing thing is that we do not necessarily perceive the room to be darker. Once we perceive that the light source is self-emitting, our brains adjust our perception of the room so we still perceive the rest of the room to be well lit. Understanding the psychophysics research related to color perception and other human visual system phenomena enable us to make more effective use of the tools that we have to create the perceptions we want to occur in our users.

How do we sense color?

Humans have (at least) four different kinds of sensors in the retina.

- One type, called the rods, sense EM radiation across a broad spectrum. Since they do not differentiate between different frequencies, rods sense in what we perceive as greyscale.
- The other type, called the cones, sense particular parts of the EM spectrum. There are three types of cone sensors that sense long, medium and short wavelengths. These correspond to what we perceive as the red, green, and blue parts of the spectrum. Different spectra produce different responses in the three types of cones. Those different patterns get converted by the brain into colors.
- There is evidence for a fourth type of sensor in the retina that is sensitive to a very small part of the spectrum that corresponds to a piece of the daylight spectrum. These sensors appear to be linked to circadian rhythms and may or may not play a role in vision or color perception.

The EM spectrum from 380nm to 780nm, which covers the range of visible light, has an infinite number of potential spectra. However, the physics of reflection and emission mean that the set of EM spectra we are likely to encounter in our lifetimes is a small subset of all possible spectra. Experiments have shown that sampling the EM spectrum and representing it with three values is sufficient to differentiate most materials and represent most spectra in the physical world.

There are, however, materials that will look identical in one lighting situation to our eyes, but different under a second lighting situation. This phenomenon is called metamerism, and it is caused by the fact we cannot sense the details of the EM spectrum.

The neurons in the human retina also have the capability to inhibit one another, so a cone of one color can end up with a negative response to certain stimuli. The relevance to computer graphics is that there are colors we can perceive—ones with negative responses in one of the color channels—that cannot be represented on a physical display device that uses additive emission to generate colors.

How do we generate color with a computer?

Color is generated on a computer monitor using some method of generating different EM spectra.

- Cathode Ray Tube (CRT): these use electron guns to activate different kinds of phosphors arranged in an array on a screen. There are generally three different colors of phosphor (red, green, and blue) and they mix together to generate different spectra.
- Liquid Crystal Display (LCD): these use small liquid crystal gates to let through (or reflect) differing amounts and colors of light. Usually, the gates are letting light pass through what are effectively R, G, or B painted glass. LCD displays are currently the most common types of displays in use.
- Plasma Display (PDP): these use very small R, G, B fluorescent lamps at each pixel. The lamps are filled with a noble gas and a trace amount of mercury. When a voltage is applied to the lamp, the mercury emits a UV photon that strikes a phosphorescent material that emits visible light. Plasma colors match CRT colors, because both use the same phosphorescent materials.
- Organic Light Emitting Diodes [OLED]: these are materials that emit photons of varying wavelengths when a voltage is applied in the correct manner. Like plasma displays, OLEDs are an emissive technology that uses collections of small OLEDs to generate a spectrum of colors.

How do we represent colors with a computer?

- The internal representation of colors matches the display (and viewing) mechanisms: mix three intensity values, one each for red, green, and blue.
- The RGB system is additive: with all three colors at maximum intensity you get white.
- There are many other color spaces we can use to represent color
 - HSI: hue, saturation, intensity
 - YIQ: intensity (Y) and two color channels (in-phase and quadrature), the original US broadcast TV standard
 - YUV: intensity (Y) and two color channels (UV), used by PAL, the original European broadcast standard
- We need to use a certain number of bits per color channel to generate different colors. Common formats include: 1, 5/6, 8, and 16 bits/pixel. High end cameras now capture in up to 14 bits/pixel.
- Internally, we will likely be using floating point data to represent pixels.

How do we represent pictures as files?

- Somehow we have to store the information in the image, which may include colors, but may also include shapes, text, or other geometric entities.
- Raster image format: each pixel gets a value
 - TIFF: Tagged Image File Format, non-lossy, any data depth is possible
 - GIF: Graphics Interchange Format, lossy, at most 256 colors, supports animation
 - JPG: long name, lossy method (can't recreate the original data), compacts images well
 - PNG: Portable Network Graphics, non-lossy compression, retains original data
 - PPM: Portable Pixel Map, very simple image representation with no compression
 - RLE: Run-Length Encoded, simple, non-lossy (but not very good) compression
- Vector image representations: images are collections of objects with geometric definitions
 - SVG: Scalable Vector Graphics, created for the web, based on XML
 - GXL: Graphics eXchange Language, also created for the web and based XML
- Element list representations: images are collections of objects and pictures
 - PICT
 - PDF
- Graphical language representations: Postscript

1.3 Getting Started...

We're going to use the PPM library for reading/writing/storing images. There are convenient tools for converting PPM images into any other format you wish, so the fact that they are not compressed or readable by a browser is not a big deal. We'll use C/C++ to write graphics programs because it's fast.

1.3.1 C/C++ Basics

C is a language not too far removed from assembly. It looks like Java (or Java was based on C syntax) but it's not. The key difference between Java and C/C++ is that in C you have to manage memory yourself. In C you have access to any memory location, you can do math on memory locations, you can allocate memory, and you can free memory. This gives you tremendous power, but remember the Spidey rule: with great power comes great responsibility. The corollary of the Spidey rule is: with great power come great screw-ups.

C is a functional language, which means all code is organized into functions. All executable C programs must have one, and only one function called `main`. Execution of the program will begin at the start of the `main` function and terminate when it returns or exits. You can make lots of other functions and spread them around many other files, but your executable program will always start with `main`.

C++ is an object oriented language that allows you to design classes and organize methods the same way you do in Java. C++ still retains its functional roots, however, and your top level program still has to be `main`. You can't make an executable function out of a class as you can in Java.

Code Organization

There are four types of files you will create: source files, header files, libraries, and object files.

- Source files: contain C/C++ code and end with a `.c` or `.cpp` suffix (`.cc` is also used for C++).
- Header files: contain type definitions, class declarations, prototypes, extern statements, and inline code. Header files should never declare variables or incorporate real code except in C++ class definitions.
- Libraries: contain pre-compiled routines in a compact form for linking with other code.
- Object files: object files are an intermediate step between source files and executables. When you build an executable from multiple source files, using object files is a way to speed up compilation.

The main function and command line arguments

One of the most common things we like to do with executable shell functions is give them arguments. C makes this easy to do. The `main` function should always be defined as below.

```
int main(int argc, char *argv[]) {  
  
    return(0);  
}
```

The `main` function returns an `int` (0 for successful completion) and takes two arguments: `argc` and `argv`. The argument `argc` tells you how many strings are on the command line, including the name of the program itself. The argument `argv` is an array of character arrays (strings). Each separate string on the command line is one of the entries in `argv`. Between the two arguments you know how many strings were on the command line and what they were.

Data Types

The basic C data types are straightforward.

- char / unsigned char: 8 bits (1 byte) holding values in the range [-128, 127] or [0, 255]
- short / unsigned short: 16-bits (2 bytes) holding values in the range [-32768, 32767] or [0, 65535]
- int / unsigned int: 32 bits (4 bytes) holding signed or unsigned integers up to about 4 billion
- long / unsigned long: 32 (4 bytes) or 64 bits (8 bytes), depending on the processor type holding very large integers
- float: 32-bit (4 byte) IEEE floating point number
- double: 64-bit (8 byte) or longer IEEE floating point number

There aren't any other basic data types. There are no native strings. You can create structures that are collections of basic data types (the data part of classes in Java). You can create union data structures where a single chunk of memory can be interpreted many different ways. You can also create arrays of any data type, including structures or unions.

```
int a[50];
float b[20];
```

The best way to create a structure is to use the typedef statement. With the typedef you can create new names for specific data types, including arrays of a particular size. The following creates a data type Vector that is an array of four floats.

```
typedef float Vector[4];
```

Example: Defining a structure

```
typedef struct {
    short a;
    int b;
    float c;
} Fred;
```

The above defines Fred to be a structure that consists of three fields a, b, and c. The syntax for accessing the fields of Fred is dot-notation. The following declares two variables of type Fred. The first is initialized in the declaration, the second is initialized using three assignment statements.

```
Fred tom = {3, 2, 1.0};
Fred f;

f.a = 6;
f.b = 3;
f.c = 2.0;
```

C does not pre-initialize variables for you (Java does). Whatever value a variable has at declaration is the result of random leftover bits sitting in memory and it has no meaning.

Strings

C does not have a built-in string type. Generally, strings are held in arrays of characters. Since an array does not know how big it is, C strings are null-terminated. That means the last character in a string must be the value 0 (not the digit 0, but the value 0). If you create a string without a terminator, something will go wrong. String constants in C, like “hello” will be null-terminated by default. But if you are manually creating a string, don’t forget to put a zero in the last place. The zero character is specified by the escape sequence ‘\0’.

Since strings in C are null-terminated, you always have to leave an extra character for the terminator. If you create a C array of 256 characters, you can put only 255 real characters in it.

Never allocate small strings. Filenames can be up to 255 characters, and pathnames to files can get large very quickly. Overwriting the end of small string arrays is one of the most common (and most difficult to debug) errors I’ve seen.

C does have a number of library functions for working with strings. Common ones include:

- `strcpy(char *dest, char *src)` - copies the source string to the destination string.
- `strcat(char *dest, char *src)` - concatenates `src` onto the end of the destination string.
- `strncpy(char *dest, char *src, size_t len)` - copies at most `len` characters from `src` into `dst`. If `src` is less than `len` characters long, the remainder of `dst` is filled with ‘\0’ characters. Otherwise, `dst` is not terminated. This is a safer function than `strcpy` because you can set `len` to the number of characters that can fit into the space allocated for `dest`.
- `strncat(char *dest, char *src, size_t count)` - appends not more than `count` characters from `src` onto the end of `dest`, and then adds a terminating ‘\0’. Set `count` appropriately so it does not overrun the end of `dest`. This is a safer function than `strcat`.

To find out about a C library function, you can always use the man pages. Typing `man strcpy`, for example, tells you all about it and related functions.

Header Files

You will want to create a number of different types for your graphics environment. In C the best way to put together new types is the typedef statement. In C++, use classes. Both types of declarations should be placed in header files. As an example, consider an Image data type. In C, we might declare the Image data type as below.

```
typedef {  
Pixel *data;  
int rows;  
int cols;  
} Image;
```

The difference with C++ and using a class is not significant, except that in C++ you can associate methods with the class.

```
class Image {
public:
    Pixel *data;
    int rows, cols;

    Image();
    Image(int r, int c);
};
```

Prototypes of functions also belong in header files. Prototypes describe the name and arguments of a function so that it is accessible in any source file and the compiler knows how to generate the code required to call the function.

```
Pixel *readPPM(int *rows, int *cols, int *colors, char *filename);
```

Extern statements are the appropriate method for advertising the existence of global variables to multiple source files. The global variable declarations themselves ought to be in source files. Initialization of the global variables also need to be in the source files. If the declaration itself is made in the header file, then multiple copies of the global variable may exist. Instead, an extern statement advertises the existence of the variable without actually instantiating it.

```
extern int myGlobalVariable;
```

Inline functions are small, often-used functions that help to speed up code by reducing the overhead of function calls. Rather than use a typical function call that requires pushing variables onto the stack, inline functions are copied into the function from which they were called. Because the functions are copied by the compiler, the compiler must have access to inline functions during compilation. Therefore, inline functions must be located in the header files. They are the only C code that belongs in a header file. In C++, methods defined within the class declaration are implicitly inline, but not necessarily. It is a good idea to only define methods explicitly declared as inline in the header file, especially for large projects.

Useful include files

Standard include files for C provide definitions and prototypes for a number of useful functions such as `printf()`, provided by `stdio.h`, `malloc`, provided by `stdlib.h`, and `strcpy()`, provided by `string.h`. In addition, all math functions such as `sqrt()` are provided by `math.h`. A good template for include files for most C programs is given below.

```
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <string.h>
```

When using C++, if you want to use functions like `printf()`, you should use the new method of including these files, given below. In addition, the include file `iostream` is probably the most commonly used include file for C++.

```
#include <cstdio>
#include <cstdlib>
#include <cmath>
#include <cstring>
#include <iostream>
```

Pointers

Variables hold data. In particular, variables hold collections of ordered bits. All variables hold nothing but bits, and what makes them different is how we interpret the bits.

In C there are two major subdivisions in how we interpret the value in a variable. Some variables hold bits that we interpret as data; it has meaning all by itself. Some variables hold addresses: they point to other locations in memory. These are pointers.

When you declare a variable, it gives a label to some memory location and by using the variable's name you access that memory location. If the variable is a simple data type (e.g. char, int, float, double) then the memory location addressed by the variable can hold enough bits for one of those types. If the variable is a pointer (e.g. char *, int *, float *, double *) then the memory location addressed by the variable can hold enough bits for the address of a memory location.

Until you allocate memory for the actual data and put the address of that allocated location into the pointer variable, the pointer variable's address is not meaningful.

You can declare a pointer variable to any data type, including types that you make up like arrays, structures and unions. Adding a * after the data type means you are declaring a pointer to a data type, not the actual data type itself. That means you have created the space for an address that will hold the location of the specified type.

To allocate space for actual data, use the malloc function, which will allocate the amount of memory request and return a pointer to (the address) of the allocated memory. The sizeof function returns the number of bytes required to hold the specified data type.

Example: Declaring and allocating pointers

```
int *a; // declare a pointer to an integer

a = malloc(sizeof(int)); // allocate memory for the integer

*a = 0; // assign a value to the integer

free(a); // free the allocated memory (the address in a is no longer valid)
```

The above first declares an int pointer and allocates space for it. The next line says to dereference the pointer (*a), which means to access the location addressed by a, and put the value 0 there. The final line frees the space allocated in the malloc statement.

Every malloc statement should be balanced by a free statement. Good coding practice is to put the free statement into your code when you make the malloc. The power of C is that you can, if you are sure about what you're doing, access, write and manage memory directly.

Arrays

Arrays in C are nothing but pointers. If you declare a variable as `int a[50];` then the variable `a` holds the address of a memory location that has space for 50 ints. The difference between `int *a;` and `int a[50];` is that the former just allocates space for the address of one (or more) integers. The latter allocates space for 50 integers and space for their address and puts the address in `a`.

The benefit to making arrays using simple pointers is that you can set their size dynamically. Because arrays are pointers, you cannot copy their value from one array to another using a simple assignment. That just copies the address of one array into the variable holding the address of the second array (which is bad). You have to copy arrays element by element.

Example: creating an array

```
int *a;
int size = 50;
int i;

a = malloc(sizeof(int) * size);

for(i=0; i<size; i++) {
    a[i] = 0;
}

free(a);
```

The above creates an array of 50 ints and puts the address of that memory in `a`. It then puts the value 0 in each of the memory locations and then frees the memory. If you want to create 500 ints, all you have to do is change the value of the variable `size`.

You can declare multi-dimensional arrays in C.

```
int a[4][4]; // creates a 4x4 array
```

However, multi-dimensional arrays in C don't act like true 2-D arrays. For a fixed-size data type, like a 4x4 matrix, they work fine. But you can't directly allocate a multi-dimensional array using `malloc`. You have to build multi-dimensional arrays yourself, as shown in the next section.

1.3.2 Creating, Reading, Modifying and Writing Images

Images are big 2D arrays. The common way to address a particular image location is (row, column) notation. This is different than traditional mathematical notation, which generally puts the horizontal axis (x-axis) first and accesses 2D locations using (x, y). The reason is that data on a screen is organized in row-major order. All of the pixels on the first row (top row) of an image come first in order from left to right. Just remember that x is a column and y is a row.

Any image can be represented as a single big array. Just take all the rows from top to bottom and concatenate them together. The index of pixel (r, c) is then $[r * \text{cols} + c]$. To allocate an image as a big array, we can just figure out how many pixels there are and allocate enough for the image.

We can also create real 2-D arrays by first creating an array of row pointers and then setting their value to the proper addresses in a pixel array. The latter makes accessing random row-column locations easier to code in many situations. Allocating, setting up, and freeing 2-D arrays takes a bit more effort. Both types of image creation and access are shown in the example below.

To read and write images we are going to use a collection of functions that read and write PPM (color) and PGM (greyscale) images. These functions are located in ppmIO.c and prototyped in ppmIO.h, which you will want to include in your programs. The function readPPM reads in an image, gets the number of rows and columns, and returns a pointer to the image data as a single array. The function writePPM writes out image data to a file as a PPM image.

```
#include <ppmIO.h>

int main(int argc, char *argv[]) {
    Pixel *image;
    int rows, cols, colors;
    int i;

    // read the image
    image = readPPM( &rows, &cols, &colors, argv[1] );

    // swap blue and red channels
    for(i=0; i<rows*cols; i++) {
        unsigned char tmp = image[i].r;
        image[i].r = image[i].b;
        image[i].b = tmp;
    }

    // write out the new image
    writePPM(image, rows, cols, colors, argv[2] );
}
```

Example: creating an image

The following creates an image as a 1-D array and initializes it to black.

```
int rows = 50;
int cols = 50;
int size = rows * cols;
int i;
Pixel *image;

image = malloc(sizeof(Pixel) * size);

for(i=0;i<size;i++) {
    image[i].r = 0;
    image[i].g = 0;
    image[i].b = 0;
}

free(image);
```

The following creates a 2-D array and initializes it to black.

```
int rows = 50;
int cols = 50;
int size = rows * cols;
int i, j;
Pixel **image; // note the image is a double pointer

image = malloc(sizeof( Pixel *) * rows ); // allocate row pointers
image[0] = malloc(sizeof(Pixel) * size); // allocate the pixels
for(int i=1;i<rows;i++) {
    image[i] = &(amp; image[0][i*cols] ); // set the row pointers
}

for(i=0;i<rows;i++) {
    for(j=0;j<cols;j++) {
        image[i][j].r = 0;
        image[i][j].g = 0;
        image[i][j].b = 0;
    }
}

free(image[0]); // free the pixel data
free(image); // free the row pointer data
```

1.3.3 Compilation Basics

C/C++ files must be compiled before you can run your program. We will be using gcc/g++ as our compiler. As noted above, you can distribute your functions/classes among many files, but only one of the files you compile together into an executable can have a main function. It's good coding style to distribute functions across a number of files in order to keep things modular and organized. As in Java, it's common to use a separate file for each class.

You can always list all of the C/C++ files you want to compile together on the command line and tell gcc to build your executable. The example below compiles the three C files and links them together into a single executable called myprogram (the -o flag tells gcc what to call the output).

```
gcc -o myprogram file1.c file2.c file3.c
```

The problem with this approach is that every time you make a change to one of the files, all of the files get recompiled. With a large project, recompiling all of the files can be time consuming. The solution to this is to create an intermediate file type called an object file, or .o file. An object file is a precompiled version of the code ready to be linked together with other object files to create an executable. The -c flag for gcc/g++ tells it to create an object file instead of an executable. An object file is always indicated by a .o suffix. To build the same program as above, you would need to precompile each of the C files and then link the object files.

```
gcc -c file1.c
gcc -c file2.c
gcc -c file3.c
gcc -o myprogram file1.o file2.o file3.o
```

This seems like a lot of extra work, until you make a change to file1.c and then want to recompile. You can rebuild the executable using the commands

```
gcc -c file1.c
gcc -o myprogram file1.o file2.o file3.o
```

This can save you lots of time when you are working with a large project, because the files you didn't touch don't have to be recompiled. gcc/g++ just uses the object file to link things together.

Sometimes you also want to link in libraries such as the standard math library or libraries that you built yourself. The -l flag lets you link in existing libraries that are in the compiler's search path. For example, to link in the math library, you would change the last line of the above example to be the following.

```
gcc -o myprogram file1.o file2.o file3.o -lm
```

When you want to link in a library you've built yourself, you need to not only link it with the -l flag, but also tell the compiler where to look with the -L flag. The example below tells the compiler to go up one directory and down into a subdirectory named lib to look for libraries. If the file libmylib.a is located there, it will successfully link.

```
gcc -o myprogram file1.o file2.o file3.o -L../lib -lmylib -lm
```

You may also need to tell the compiler where to find include files (.h files). If, for example, all of your include files in an include subdirectory, you will want to add a -I (dash capital I) option to the compile line when compiling the .c files. Linking is too late for includes, as they must be available during compilation. The following compiles the file file1.c into an object file file1.o and tells the compiler to look in a neighboring include subdirectory called include.

```
gcc -c file1.c -I../include
```

There are many other flags that can be used by gcc. Two you will see this semester are `-Wall`, which turns on all warnings, and `-O2`, which sets the optimization level to 2, which is fairly aggressive. It will likely make your code a bit faster, which is good.

Development Tools

The mac does not come with development tools in a standard install, although they are included on the installation disks. You need the Xcode package to get gcc/g++ and standard development tools like make. If you also want X-windows and other Unix tools and programs, then you will want to download XQuartz (version 2.3 as of September 2008) and MacPorts. MacPorts is a package manager and lets you install things like ImageMagick, gimp, and xv. Install Xcode, the XQuartz, the MacPorts. It will take a while, but none of it is particularly difficult.

Windows also does not come with development tools standard. The cygwin package is the most common way developers get an X-windows environment and all of the standard unix development tools. Note that there may be additional packages you need to download in addition to the standard install.

1.3.4 Makefiles

The make program is a way of automating the build process for programs. It's most useful for large programs with many files, but it's also useful for smaller programs. It can save you from typing a lot of stuff on the command line and keeps you from having to remember all the flags every time.

The simplest makefile is just two lines. The first line defines the name of the rule and its dependencies, the second line defines the rule's action.

```
myprogram: file1.c file2.c file3.c
    gcc -o myprogram file1.c file2.c file3.c -I../include -lm
```

The first line defines a rule called `myprogram` and lists the files upon which the rule depends. If any of those files change, the rule should be executed. The second line defines the action of the rule. If you type `make myprogram` the rule will execute. In fact, since make executes the first rule by default, all you need to type is `make` and the rule will execute.

When you're just starting out, create simple makefiles where the rules are spelled out explicitly. There are many more things you can do with make, and it incorporates a complete scripting language for automating complex tasks. For more information, see the makefile tutorial linked to the course website.

1.4 The Art of Graphics

Graphics is ultimately about creating images. A good artist knows about their canvas, their paints, and their brushes and uses the appropriate tools depending upon the scene they want to create. The same is true for computer graphics. The more you understand about your canvas, paints, and brushes the better images you can create.

Unlike physical artists, a digital artist is not limited in their vision of a brush. Any digital tool can be used as a building block to generate imagery, including hierarchical combinations of tools and procedural or algorithmic tools.

1.4.1 Canvas

Digital images or image sequences are most easily displayed electronically on a 2D raster screen such as a CRT, LCD, or plasma display.

- Raster images are collections of discrete pixels.
- Electronic displays use an additive color system (RGB).
- Each pixel contains all the data required to display itself.
- The information represented in the image is constant; each refresh requires a read through the entire **frame buffer**—the memory in which the image to be displayed is stored.
- Each pixel represents a finite area that will be a single color, which can make it difficult to render high frequency details without high spatial resolution. The spatial resolution of the screen, and the shape of the pixel itself creates a situation where we need to think about how to render high spatial frequency information on a lower spatial frequency device. The most common artifacts in computer graphics—the **jaggies**, or stair-step artifacts on what should be smooth continuous shapes—come from this property of raster displays.
- The retina display—a raster display where the pixels are so small that normal human vision cannot distinguish them—creates a fundamental change in the way we think about rendering graphics. When the display is at a higher resolution than the material we need to display, it removes the need to worry about artifacts like jaggies.
- The **display refresh rate** is an important value to know. The image is updated at a rate between 60-75Hz (Hz = 1/s). Therefore, when creating real time animation, your system has between 13-16ms to generate the next frame.
- The system generally has two **frame buffers**: a frame buffer is the memory in which the image to display is stored. Normally, one frame buffer is being transferred to the display while the other frame buffer is being written by the graphics engine.

Vector display systems, which are rare today, provide a different kind of canvas.

- Vector displays have a single electron gun and single color long-glow phosphor. The phosphorous on the screen has an extremely fine grain size, meaning the spatial resolution is at or beyond the limit of our eyes to distinguish individual elements.
- Vector displays generally draw lists of lines represented as point pairs.

- For each line, the electron gun moves from the starting point to the ending point, drawing a line that is the width of the electron beam.
- The system draws exact lines between high resolution locations; the coordinate system does not need to be integral, but exists at the resolution of the electron gun control system.
- The speed of the system determines the maximum number of lines that can be drawn per second. Therefore, the more lines there are in the image, the longer it takes to draw the entire list of lines. The speed of animation sequences, therefore, is determined by the complexity of the scene.
- The only parts of the screen affected are the lines, and they are drawn at the resolution of the electron gun focus and phosphor grain.

Hard-copy devices such as laser printers, inkjet printers, and plotters are common ways to view images.

- Most hard copy devices are raster devices with a high resolution, but not necessarily high enough to avoid spatial artifacts.
- Plotters that worked like vector displays used to be more common. They take longer to print, but can produce artifact-free copy.
- Hard-copy devices must use a subtractive system (CMYK) where each kind of ink subtracts one band (R, G, or B).
- Synchronizing printer output and an electronic display is non-trivial. People who depend upon high quality output use calibration equipment to guarantee that the color on the display is the color that comes out of the printer.
- 3-D printing devices are becoming more common, enabling the physical realization of 3D models.

1.4.2 Paints

While the internal representation of an image is almost always RGB (3-channel) or greyscale (1-channel), there are other ways to describe colors.

- RGB—red, green, and blue—is an additive color system and is the most common internal representation of colors. It roughly matches the output display device representation, so the complexity of the transformation of color information from internal representation to display is minimized. RGB color space is usually visualized as a cube. While RGB is useful as a data representation, it is not intuitive as a color selection space (e.g. what is orange?).
- HSI/HSV—hue, saturation, and intensity or value—is usually visualized as a cylinder. The intensity axis is the cylinder's central axis, saturation is distance from the central axis, and hue is the angle around the circle. HSI is a more intuitive color space for people because intensity and color are more clearly separated and the hue captures what we generally think of as color names (red, orange, yellow, green, blue, indigo, violet).
- Ad hoc representations: think about color picker options and all the various ways colors can be represented (e.g. crayons).
- High spectral resolution representations: maintain continuous or high frequency discrete representations of spectra throughout the rendering process.

1.4.3 Brushes

Input devices have always brought out tremendous creativity in engineers and scientists. In computer graphics, a brush is any device that can capture information from the world or generate it synthetically. That is a pretty large set of tools.

- Buttons (keyboard)
- Dials
- Position measurement (mice, trackpads, tablets, data-gloves, body suits, mo-cap systems)
- 2-D scanners
- 3-D scanners
- Motion capture
- Digital images and videos
- Stereo imagery
- Satellite/radar data
- Haptic inputs
- Microphones/music

Manual paint tools (e.g. Canvas, Photoshop, Gimp)

- Create 2D images
- Create 3D images
- Modify existing pictures
- Create textures

Plotting pixels (write code or use MacPaint)

- Tough to create sophisticated images this way
- Pixels are the most basic unit of creativity in graphics, so an artist will almost always want access to individual pixels to create a finished product.

Graphics Primitives

- Lines, curves, rectangles, ellipses, splines
- One level of abstraction above pixels
- Filling shapes affects large areas of pixels

3D Models

- 3D lines, spheres, cylinders, prisms, polygons, blobs
- Teapots, trees, buildings, other shapes in a library
- Free-form surfaces: Bezier surfaces, NURBS, subdivision-surfaces
- Another level of abstraction above graphics primitives
- Permits design in a 3D space, but requires a path to view the space

Procedural Models

- Landscape brushes, vegetation models
- Building cities from algorithms
- Weathering stone
- Mountains
- Builds upon simple 3D models to create complex aggregate models

Renderer

- Converts 3D models into a 2D image
- Incorporates models of appearance
- Requires designing the viewpoint and the view
- The Director's job

Animation

- Many levels of abstraction
- Manually specify points
- Manually specify key frames and let the computer interpolate the rest
- Manually specify paths (3D curves)
- Manually specify joint relationships
- Procedural animation - write scripts that automate the above
- AI - write scripts that guide agents and their behavior

Non-photorealistic Rendering

- Forget realism, go for artistic styles
- Has practical application in technical drawings
- Abstracts away from pixels and standard graphics primitives
- Some NPR methods are hyperrealistic imitations of physical artistic forms

2 Graphics Primitives

Graphics primitives are the next level of abstraction above pixels. They are visual representations of geometric objects such as squares, rectangles, circles, ellipses, and so on. It is important to remember that they are representations of such objects; a raster system cannot display geometric primitives exactly

2.1 Pixels

Somehow we need to tell an image how to set the color of a pixel.

- `Image_set(?)`
 - Needs access to an Image
 - Which pixel? One index or two?
 - What value?

The above raise questions about how we're going to organize our code and give functions the information they need to do their job. One extreme is to make the graphics system **state based**. The other extreme is to make the system **parameter based**.

A state based graphics system maintains a set of internal state variables. State variables might include the following attributes.

- Image
- Color
- Pen width
- Pen style or pattern
- Whether to fill graphics primitives

The graphics system will begin in a default state and the programmer will have access to functions that get or set the state variables. To draw a single red point with a 2x2 size, for example, the programmer would make three function calls.

1. Set the color to red
2. Set the pen width to 2
3. Plot the point using the current graphics state

A parameter based graphics system does not maintain state variables, but relies on the programmer to keep track of all the graphics system information and pass it in as arguments to functions that draw primitives. To draw the same red point in a 2x2 size takes only a single function call in a parameter based system, but the function has to take many more arguments.

There are strengths and weaknesses to both approaches. A state based system can be cumbersome when drawing many unique objects, requiring multiple function calls for each graphics primitive. A parameter based system, however, is more cumbersome when drawing many objects that have the same attributes. Clearly, there are compromise systems that combine aspects of both. A simple improvement on a parameter

based system, for example, is to collect all attributes in a single structure and just pass along a pointer to the structure every time a graphics primitive function is called.

There are commercial examples of both kinds of graphics systems. X-windows, for example, uses a parameter based system where the graphics attributes are bundled into a data structure passed along from function to function. OpenGL, on the other hand, is a state-based system where most graphics attributes are set by modifying state variables. For this class we're going to take a more parameter based approach to building a graphics system so that the attributes are more transparent and we avoid the need for global variables (for the most part).

Plotting pixels in our system is going to require that we pass in all the required information necessary to execute the operation. The first design decision we have to make is how to internally represent pixels. While many image formats use unsigned characters (8-bit integers) to represent each color, the calculations we will be making based on colors, lighting and geometry will require floating point operations and make more sense on range of [0, 1]. Therefore, internally, we will want to use a floating point value for each color channel.

In addition, we may want our pixel type to incorporate an alpha channel and a depth channel. The latter will be necessary once we begin to do 3-D hidden surface removal in our rendering system. How you organize your pixel internally does not make much difference, although using an array to represent the colors may make certain operations possible using a loop structure. Something like the following would work.

```
typedef struct {
    float rgb[3];
    float a;
    float z;
} FPixel;
```

An internal image, therefore, will consist of an array of FPixels and basic information like the number of rows and columns in the image.

```
typedef struct {
    long rows;
    long cols;
    FPixel *data;
} Image;
```

With these factors in mind, we might define the following basic functions to access pixels.

- FPixel image_get(Image *src, int r, int c);
- FPixel image_get1D(Image *src, int i);
- void image_set(Image *src, FPixel p, int r, int c);
- void image_set1D(Image *src, FPixel p, int i);

2.2 Graphics Specification Document

Each assignment this semester will be accompanied by a specification of the functions required for the assignment and the names and arguments for those functions. If your code follows the specifications, then you will be able to execute the test function(s) provided for each assignment. It also guarantees that you will be able to continue to build upon your code from week to week. While there will be some situations where you have to go back and add new features to old code, following the specification should make this straightforward.

Note that you have great latitude to go beyond the specification. The specification is a minimalist document and outlines the minimum required capabilities for your system. You are encouraged to add additional functionality throughout the semester. As long as the functions outlined in the specification exist and work properly, the test programs will run.

2.3 Lines

How do we draw lines into an image?

Line equation:

$$y = mx + b \quad (1)$$

Given a starting and end point, we can solve for the slope m and intercept b

$$m = \frac{y_1 - y_0}{x_1 - x_0} \quad (2)$$

$$b = y_0 - mx_0 \quad (3)$$

The differential equation for a line tells us how much to move in y given a unit step in x .

$$\Delta y = m \Delta x \quad (4)$$

To draw a line, we could start at one end (assume the left end, always), draw the first pixel and then take a step, updating x or y appropriately.

$$y_{k+1} = y_k + m \quad (5)$$

$$x_{k+1} = x_k + \frac{1}{m} \quad (6)$$

If the slope is less than one, use x as the loop variable. If we increment x by 1 each time, then we increment y by the slope m . If the slope is greater than 1, use y as the loop variable and increment x by $\frac{1}{m}$ each step.

How do we have to represent x and y ? Given fractional values, which pixel do we light up? On older hardware, using floating point values was completely out of the question as it was 100s of times slower than integer math. Even on new hardware, if we do line calculations using integers we save hundreds of thousands of transistors on a graphics card that can be used for other tasks.

2.3.1 Bresenham's line algorithm

Bresenham's algorithm is based on analyzing how an error term changes as you move from one pixel to the next. Figure 1 shows the distances d_1 and d_2 , which represent the distances from the line to the upper and lower pixel centers, respectively.




Figure 1: Distance of a line to upper pixel center d_1 and lower pixel center d_2 .

If we define an error term $e = d_2 - d_1$, we can look at the sign of e to determine which pixel to light up.

- If the sign is positive, light the upper pixel.
- If the sign is negative, light the lower pixel.

Each iteration of the algorithm we step to the next pixel along the x axis, update the error term $e = e + m$ and look at the new sign of e . We never actually have to calculate $d_2 - d_1$, explicitly.

- If e is positive, then we take a step up and subtract 1 from e to account for the step.
- Initialize $e = -\frac{1}{2}$ so the first decision will be correct (consider the case of a $\frac{1}{2}$ slope).
- The first value of e we test will be $e = -\frac{1}{2} + m$

It's still not an integer algorithm!

Integer Bresenham's algorithm

A few simple modifications to the algorithm can convert it to an integer algorithm.

- Define $e' = 2\Delta x e$, so the initial value of $e' = 2\Delta x(-\frac{1}{2}) = -\Delta x$
- Represent $m = \frac{\Delta y}{\Delta x}$, both Δy and Δx are integer values, and $2\Delta x(m) = 2\Delta x\left(\frac{\Delta y}{\Delta x}\right) = 2\Delta y$
- Consider the first value of e' we test

$$e' = 2\Delta x\left(-\frac{1}{2} + m\right) = 2\Delta x\left(-\frac{1}{2} + \frac{\Delta y}{\Delta x}\right) = 2\Delta y - \Delta x \quad (7)$$

- The update rule for e' is $e'_{k+1} = e'_k + 2\Delta y$
- When $e' > 0$ we subtract $2\Delta x$.

Pseudo-code algorithm (for lines in first octant going left to right)

```

x = x0
y = y0
dx = x1 - x0
dy = y1 - y0
e = 2 * dy - dx

for(i = 0; i<dx; i++) {
    Image_set(im, color, y, x)
    while( e > 0) {
        y = y + 1
        e = e - 2 * dx
    }
    x = x + 1
    e = e + 2 * dy
}

```

Example: Drawing a line

Let $(x_0, y_0) = (20, 10)$ and $(x_1, y_1) = (30, 18)$, so $(\Delta x, \Delta y) = (10, 8)$.

i	draw pixel	e_{k+1}	x_{k+1}	y_{k+1}
Initial	n/a	$6 = 2 * 8 - 10$	20	10
0	(20, 10)	2	21	11
1	(21, 11)	-2	22	12
2	(22, 12)	14	23	12
3	(23, 12)	10	24	13
4	(24, 13)	6	25	14
5	(25, 14)	2	26	15
6	(26, 15)	-2	27	16
7	(27, 16)	14	28	16
8	(28, 16)	10	29	17
9	(29, 17)	6	30	18

To generalize Bresenham's algorithm always draw the line from bottom to top and consider the cases of the first four octants.

- First octant uses the algorithm as specified above.
- Second octant flips the roles of x and y.
- Third octant is like the second, but steps left in x instead of right.
- Fourth octant is like the first, but steps left in x instead of right.
- Horizontal and vertical lines must be handled as special cases.

Faster line algorithms

- Can take advantage of symmetry and draw in from both ends simultaneously
- Can take multiple steps at a time. There are only four possible arrangements of 2 pixels in a line.

2.4 Screen coordinate systems

Lines, polygons, circles, and ellipses are mathematic objects with definitions in a continuous world. In the absence of a vector display, we have to render these continuous objects on a discrete raster display. When we display graphics primitives, we would like to have attributes like length and area be the same in their rendered state as they are in their theoretical state.

Consider the example above. If we consider the coordinate grid to be located at pixel centers, then we probably want to light up both the first and last pixel on the line: (20, 10) and (30, 18). However, if we do so, the line is no longer the correct length in its visible representation. If we stop one before the last pixel, the line is the correct length, but the choice of whether to light up the first or last pixel seems arbitrary.

What if we consider the coordinate system to correspond to the left left corner of a pixel?

- Makes sense to not color in the last pixel
- Need to be careful how we draw lines
- Need to think about any impact this has on Bresenham's line algorithm

2.4.1 Updating Bresenham's algorithm

If we move the coordinate system to the lower left corner, then the line slides down and left by half a pixel. The decision to move up or not is still the same: look at whether the line is above or below the pixel boundary at the halfway point.

- Original initial testing point: $e'_0 = 2\Delta x(\frac{\Delta y}{\Delta x} - \frac{1}{2}) = 2\Delta y - \Delta x$
- Modified initial testing point is down by half a pixel $-\frac{1}{2}$ and left by half a pixel $+\frac{1}{2}\frac{\Delta y}{\Delta x}$:

$$e'_0 = 2\Delta x \left(\frac{\Delta y}{\Delta x} - \frac{1}{2} + \left(\frac{\Delta y}{2\Delta x} - \frac{1}{2} \right) \right) = 2\Delta x \left(\frac{3\Delta y}{2\Delta x} - 1 \right) = 3\Delta y - 2\Delta x \quad (8)$$

Since the increments to the error are based on unit steps, and we're still doing unit steps, the incremental values do not change when using a lower-left corner origin for the coordinate system.

How does this affect the lines? Consider where the cutoff points are for stepping up or across.

- $[0, \frac{2}{3}]$: step horizontally
- $(\frac{2}{3}, \frac{3}{2})$: step diagonally
- $[\frac{3}{2}, \infty]$: step up

What about horizontal and vertical lines? The concept of line direction becomes important, especially when drawing rectangles. One convention is to make horizontal and vertical lines rotated versions of one another.

- Draw horizontal lines going right in the first quadrant
- Draw vertical lines going up in the second quadrant
- Draw horizontal lines going left in the third quadrant
- Draw vertical lines going down in the fourth quadrant

If you draw closed polygons using horizontal and vertical lines in a counterclockwise order using the convention above, then the area enclosed by the lines is mathematically correct. Drawing the same set of points in the opposite order results in missing corners. If you try to draw all vertical lines in the first quadrant and all horizontal lines in the first quadrant, then there isn't any ordering of the corners that produces a mathematically correct solution: you have to draw to different points.

Programs like Adobe Photoshop take into account where the user clicks in the image when drawing lines. If you are zoomed in, Photoshop actually adapts the starting location of the line to a subpixel location. Clicking near the lower left corner of a pixel produces a line that is identical to one produced by Bresenham's with a lower left starting point.

2.5 Polygons and polylines

A **Polyline** is simply a collection of connected line segments where the end of one line is the beginning of the next. Polylines do not connect the last point to the first. A polyline function will take in a list of points and graphics attributes and draw the line segments.

A **Polygon** is a collection of connected line segments where the last point is connected back to the first point by a line. A polygon function takes the same arguments as the polyline function and just adds the last connecting line segment.

2.6 Circles and Ellipses

How can we draw a circle into an image?

Algorithm 1: cartesian circle equation

- Circle equation: $r^2 = (x - x_c)^2 + (y - y_c)^2$
- Solving for y

$$y = y_c \pm \sqrt{r^2 - (x_c - x)^2} \quad (9)$$

- Step across in x , calculating the next y .
- Switch orientations half-way through.

Problems:

- Non-uniform spacing
- Get points along the circle, not a complete, closed shape
- Square roots are expensive

Algorithm 2: polar circle equation

We can improve the spacing by using an angular formulation of a circle.

$$\begin{aligned} x &= x_c + r \cos(\theta) \\ y &= y_c + r \sin(\theta) \end{aligned} \quad (10)$$

If you walk around the circle by modifying (θ) in steps of $\frac{1}{r}$ you can get unit steps. However, it makes use of trigonometric functions, which makes it costly.

Algorithm 3: Bresenham's Circle Algorithm

The main idea of the integer circle algorithm is to use the cartesian circle function to test whether to step across or down by evaluating the relative position of the two possible next pixels to the circle boundary. We can do this in a single test by evaluating the circle equation on a point that is mid-way between the pixels.

- $f(x, y) = x^2 + y^2 - r^2$
- If (x, y) is our test point, then the sign of $f(x, y)$ tells us if the boundary is inside the circle (< 0) on the circle ($= 0$) or outside the circle (> 0).

We actually know the first point to draw. Relative to the origin (x_c, y_c) , the first point is one of the four pixels horizontally or vertically from the center. Now we just need to calculate the next location to test. Ideally, we'd like to just keep track of an error term, update the error term with each step, and choose which pixel to draw next based on the error term. Note that circles are radially symmetric, so we only need to figure out the points in one octant.

Derivation:

- The last point drawn is (x_k, y_k)
- The next test location is $(x_k + 1, y_k - 1/2)$

- The next value to test is

$$p_k = f(x_k + 1, y_k - 1/2) = (x_k + 1)^2 + (y_k - 1/2)^2 - r^2$$

$$p_k = x_k^2 + 2x_k + 1 + y_k^2 + y_k + 1/4 - r^2$$

$$p_k = x_k^2 + y_k^2 + 2x_k + y_k + 5/4 - r^2$$

- The value to test on the next iteration is

$$p_{k+1} = f(x_{k+1} + 1, y_{k+1} - 1/2) = (x_{k+1} + 1)^2 + (y_{k+1} - 1/2)^2 - r^2$$

$$p_{k+1} = x_{k+1}^2 + 2x_{k+1} + 1 + y_{k+1}^2 + y_{k+1} + 1/4 - r^2$$

$$p_{k+1} = x_{k+1}^2 + y_{k+1}^2 + 2x_{k+1} + y_{k+1} + 5/4 - r^2$$

$$p_{k+1} = (x_k + 1)^2 + y_{k+1}^2 + 2(x_k + 1) + y_{k+1} + 5/4 - r^2$$

$$p_{k+1} = x_k^2 + y_{k+1}^2 + 2(x_k + 1) + 2x_k + y_{k+1} + 1 + 5/4 - r^2$$

- Subtract p_k from p_{k+1} to get the update rule

$$p_{k+1} = p_k + 2[x_k + 1] + [y_{k+1}^2 - y_k^2] - [y_{k+1} - y_k] + 1 \quad (11)$$

- y_{k+1} is either y_k or $y_k - 1$

Consider that the second and third terms involve a difference of y values, which will be either 0 if no step was taken, or a known increment if a step was taken. How we update p depends upon whether we stepped across ($y_{k+1} = y_k$) or stepped down ($y_{k+1} = y_k - 1$). In all cases, we know that $x_{k+1} = x_k + 1$.

- If p_k is negative: $p_{k+1} = p_k + 2x_{k+1} + 1$
- if p_k is non-negative: $p_{k+1} = p_k + 2x_{k+1} + 1 - 2y_{k+1}$

Note that all of the updates are integral, so there is no need for floating point values. The only exception is that the initial starting value for p is $5/4 - r$, but rounding the expression to $1 - r$ works just fine.

The algorithm for drawing a circle, then, is to begin at one of the axis extrema and then iterate over one octant, reflecting each drawn pixel to the other eight octants. At each iteration test the value of p_k and step either across or down, updated the value of p_k appropriately given the rules above.

Ellipses

The ellipse algorithm is almost identical to the circle algorithm except that you have to draw a full quadrant of the ellipse and change the step variable at some point in the process. The quadrant divides where the slope is -1, or the first pixel where the change in x is less than the change in y . Otherwise, the process is the same: take a step, update an error term, and reflect the points across to the other four quadrants.

2.6.1 Coordinate System Issues

Like lines, we have to be careful about how to draw ellipses and circles or they end up being too big. The solution for both primitives is to make all calculations in the third quadrant and reflect the points around to the others using the diagram in figure 2. For the circle algorithm, this modifies the initial values of x and y and modifies the update rule because the x values are negated. Note that all the x and y values are relative to the circle's center, which is used only when actually plotting the pixels.

- Initial value of $x, y = -1, -r$
- Initial value of $p = 1 - r$
- Decrement x each iteration
- Increment y each iteration (moving up and left in the third quadrant)
- Iterate until $x > y$
- Update rule if $p < 0$: $p_{k+1} = p_k + 1 - 2x_{k+1}$, where $x_{k+1} = x_k - 1$
- Update rule if $p \geq 0$: $p_{k+1} = p_k + 1 - 2(x_{k+1} - y_{k+1})$ where $y_{k+1} = y_k + 1$ and $x_{k+1} = x_k - 1$
- When reflecting, subtract one from the reflected value any time the x or y coordinate is negated




Figure 2: Reflecting points around the octants

Corrected Circle Algorithm

```
Parameters: cy, cx, radius
set x to -1
set y to -radius (rounded)
set p to 1 - radius (rounded)
Plot (x + cx, y + cy) and reflect it around the circle
```

```
while x > y
  x--
  if p < 0
    p = p + 1 - 2 * x
  else
    y++
    p = p + 1 - 2 * (x - y)
  plot ( x + cx, y + cy ) and its reflections
```

For ellipses, you get the following modifications, which are meaningful in the context of the example code provided with the lab.

- Initial value of $x = -1$
- Initial value of $y = -R_y$ (y-axis radius)
- Initial value of $p_x = 2R_y^2$
- Adjust the initial value of p in section 1 by adding $R_y^2 + p_x$
- Adjust the initial value of p in section 2 by adding $R_x^2 - p_y$
- As with the circle, when reflecting around, subtract 1 from a coordinate whenever it is negated.

Corrected Ellipse Algorithm

Parameters: cx , cy , R_x , R_y

Initial value of $x = -1$

Initial value of $y = -R_y$

Initial value of $p_x = 2 * R_y * R_y$

Initial value of $p_y = 2 * R_x * R_x * -y$

Plot $(x + cx, y + cy)$ and its four reflections

$p = R_y * R_y - R_x * R_x * R_y + R_x * R_x / 4 + R_y * R_y + p_x$

while $p_x < p_y$

$x--$

$p_x = p_x + 2 * R_y * R_y$

 if $p < 0$

$p = p + R_y * R_y + p_x$

 else

$y++$

$p_y = p_y - 2 * R_x * R_x$

$p = p + R_y * R_y + p_x - p_y$

 Plot $(x + cx, y + cy)$ and its four reflections

$p = R_y * R_y * (x * x + x) + R_x * R_x * (y * y - 2 * y + 1) - R_x * R_x * R_y * R_y + R_x * R_x - p_y$

while $y < 0$

$y++$

$p_y = p_y - 2 * R_x * R_x$

 if $p > 0$

$p = p + R_x * R_x - p_y$

 else

$x--$

$p_x = p_x + 2 * R_y * R_y$

$p = p + R_x * R_x - p_y + p_x$

 Plot $(x + cx, y + cy)$ and its four reflections

2.6.2 Filled Circles and Ellipses

Fill circles and ellipses is trivial

- For each horizontal pair of points, fill in the row from left to right.
- For each point you calculate for a circle, you get four lines
- For each point you calculate for an ellipse, you get two lines

2.6.3 Why Primitives are Important

Trying to draw simple graphics primitives quickly and correctly raises many issues that are essential to graphics.

1. Fast algorithms use differential methods and avoid computing (x, y) from scratch
2. Deriving differential algorithms involves the following process
 - identify the initial point
 - identify a test location
 - calculate the error at iteration i
 - calculate the error at iteration $i + 1$ in terms of the variables at step i
 - subtract the two error terms to identify the update rule for the error value
3. Integer algorithms exist for most basic primitives
4. Picking the right coordinate frame is important
 - impacts the geometric qualities of the displayed object (area, perimeter, length)
5. Trying to render continuous mathematical objects on a discrete screen results in aliasing artifacts

2.7 Generic Fill Algorithms

Often we want to fill arbitrary shapes or polygons. The two major approaches are flood-fill and scan-fill. Flood-fill algorithms start at a seed point inside the shape or area and grow out from the seed until they hit a boundary. Scan-fill algorithms begin at the top of the shape and work down row by row, figuring out which pixels to color on each row. Flood fill algorithms are quick and easy to implement and work well for interacting drawing programs. Scan-fill algorithms are more appropriate for automatic rendering. Scan-fill algorithms also provide a framework for doing many more complex tasks while drawing the shape.

2.7.1 Flood-fill Algorithms

The **flood fill** algorithm is implemented as the paint can tool in most drawing programs. The idea is to fill in an area surrounded by a boundary, starting at a seed pixel and growing outwards. The flood-fill algorithm is straightforward to implement using a stack.

```
Make the boundary using a selected color C
Pick an initial point
Color the initial point
Push the initial point on a stack
```

```
While the stack is not empty
  Pop a pixel P off the stack
  For each neighbor N of P
    If N is not already colored
      Color N
      Push N on the stack
```

There are a number of variations to the algorithm.

- If you want the boundary to be a different color than the interior, modify the if-statement so it doesn't color the pixel N if its existing color is the boundary color or the fill color.
- If you want the new color to fill over only one existing color C_{old} , modify the if-statement so that it colors pixel N only if the existing color is C_{old}

A variation on the flood-fill algorithm is the **scanline flood fill** algorithm that is a bit faster and generally uses less memory. The idea behind the scanline flood-fill is to fill in all of the contiguous pixels on a row before moving up or down to the next row.

- Fill in the scanline from left to right
- At each pixel, look at the row above and the row below to see if a seed needs to be pushed onto the stack
- Seeds need to be placed any time the scan hits the end of a contiguous span of pixels to be colored

Figure 3 shows the sequence of seeds for a polygonal shape with some concave boundaries.

Algorithm

```

Make the boundary using a selected color C
Pick an initial point (uppermost, leftmost pixel to be colored)
Color the initial point
Push the initial point on the stack

While the stack is not empty
  Pop a pixel P off the stack
  Find the leftmost pixel in the scanline that should be colored
  For each pixel Q from the leftmost to the rightmost fill location on the scanline
    Color Q
    If the pixel above Q should be colored and
    the pixel up and right from Q is on the boundary
      Push the pixel above Q onto the stack
    If the pixel below Q should be colored and
    the pixel down and right from Q is on the boundary
      Push the pixel below Q onto the stack
  If the pixel below the last pixel on the scanline should be colored
    Push the pixel below onto the stack
  If the pixel above the last pixel on the scanline should be colored
    Push the pixel above onto the stack

```




Figure 3: Diagram of scanline flood fill showing all of the seeds generated in the process. The initial seed is #1. Note the two #6 seeds because the scanline is broken by the interior boundary

2.7.2 Connectedness

We want outlines on a raster display to form a connected series of pixels. But the definition of connectedness is not always clear. Are only NSEW neighbors connected, or are diagonal pixels also connected?

4-connected: only the NSEW neighbors of a pixel are considered to be connected

8-connected: all eight neighbors of a pixel are considered to be connected

Note that the connectedness of an outline influences the definition of neighbors while filling. If a boundary is 4-connected, then you want to fill using an 8-connected definition. If a boundary is 8-connected, then you want to fill using a 4-connected definition. If you don't use a fill connectedness that is the opposite from the boundary connectedness, then the fill either will not be complete or will bleed through the boundary.

2.7.3 Beyond Single Colors

There are many ways to draw a line

- Single color
- Gradient
- Dashed line
- Transparent line

Likewise, there are many different ways to fill a shape.

- Single color
- Gradient fill
- Pattern fill
- Transparent fill
- Texture fill

How can we implement these different methods of rendering shapes?

Gradient fill

To implement a gradient fill we need two colors— C_a and C_b —and a method of choosing, for each pixel, what percentage of each color to use. Then we can use alpha blending to create the color to draw onto the screen.

$$C_s = \alpha C_a + (1.0 - \alpha) C_b \quad (12)$$

For drawing a line that blends evenly from one color to the other, we just need to keep track of the change in α as we draw the line. If the initial value of α is 0, then the final value should be 1. In the first octant, the number of steps is given by dx , so the update equation for alpha is given by the following.

$$\alpha_{i+1} = \alpha_i + \frac{1}{dx} \quad (13)$$

It's possible to execute alpha blending in a completely integral fashion if you multiply α by dx and adjust the blending equation accordingly.

For drawing a shape like a circle or ellipse that blends evenly from one color to another, there are many more options. The simplest is to pick a horizontal or vertical orientation for the gradient. Then α begins at 0 on one side of the shape and goes to 1 on the other side. However, there are many other possible blending strategies. For example, you may want the center of the circle to be color C_b and the outside to be color C_a . In that case, α is related to the radial distance of the pixel from the center of the circle.

$$\alpha_{x,y} = \frac{\sqrt{(x - x_c)^2 + (y - y_c)^2}}{r} \quad (14)$$

In fact, you can make any relationship you wish between α and location within the shape, including picking α from a random distribution. The key is to develop a set of parameters for drawing gradient fill that gives the programmer/user flexibility without making the function unwieldy.

Pattern fill

Most graphics programs give you the option of selecting a repeating pattern or texture for drawing filled objects. The patterns are generally modifications of the foreground color (or you can think of them as patterns of alpha-blending). The textures are generally designed so they flow seamlessly from one tile to the next in any direction.

At each pixel, we just need to know what part of the pattern to draw. There is a one-to-one correspondence between the patterns or textures and pixels, so there is no sampling issue, we just need to know which pixel to grab from the pattern.

The simplest method of picking a pattern pixel is to use modulo arithmetic. If the pattern is 16×16 , then if you take the coordinates of the pixel you want to draw modulo 16, you get an index into the pattern. The only issue with using the raw coordinate is that if you draw the graphics primitive in a different location, the pattern may be different with respect to the object's origin. As the object moves across the screen, the pattern stays constant with respect to the image coordinate system.

A simple modification solves the problem. Simply subtract off the object's origin (often the upper left-most pixel of the bounding box) from the pixel coordinates of interest before executing the modulo operator.

$$\begin{aligned} x_{\text{pattern}} &= (x_{\text{image}} - x_{\text{origin}}) \bmod \text{PatternSize}_x \\ y_{\text{pattern}} &= (y_{\text{image}} - y_{\text{origin}}) \bmod \text{PatternSize}_y \end{aligned} \quad (15)$$

Transparent fill

To fill a shape transparently is simply a variation on alpha blending. To have a uniform transparent fill, select α and then blend the foreground fill color C_f with the existing color in the image C_i .

$$C = \alpha C_f + (1.0 - \alpha) C_i \quad (16)$$

You can use the same kinds of functions to determine α for transparent filling as you do for gradient fill. The only difference is that instead of specifying a background color, use the current image values.

Texture fill

Texture fill is a little more complicated because the texture is likely to be some arbitrary image, and the object you want to fill is not likely to have the same shape. The key to texture fill is finding a mapping from the pixel you want to draw to the pixel in the texture you want to use.

For example, if we want to fill a circle with some part of an image, we need to find a mapping from a pixel on the circle to the texture image. A simple way to execute the mapping is to enforce a 1-1 mapping from image to texture and match the bounding box of the circle, defined by $(x_c - r, y_c - r)$ and $(x_c + r, y_c + r)$ to the upper left corner of the texture image. The coordinates in the texture image (i_t, j_t) would correspond to pixel (i, j) in the circle as in (17).

$$\begin{aligned} i_t &= i - (y_c - r) \\ j_t &= j - (x_c - r) \end{aligned} \quad (17)$$

If we let the user pick an origin for the box in the texture image, then we need to add the offset of the box (t_x, t_y) in the texture image relative to the upper left corner.

$$\begin{aligned}i_t &= i - (y_c - r) + t_y \\j_t &= j - (x_c - r) + t_x\end{aligned}\tag{18}$$

What if we let the user pick any axis-oriented non-zero box to map onto the circle? The process of mapping from circle to texture image then becomes: align the coordinate systems, scale the coordinate systems, add the offset. If we let the box in the texture image be defined by the point pair (p_{0x}, p_{0y}) and (p_{1x}, p_{1y}) , then the scale factors are the texture box width and height divided by circle bounding box size $2r$.

$$\begin{aligned}s_x &= \frac{p_{1x} - p_{0x}}{2r} \\s_y &= \frac{p_{1y} - p_{0y}}{2r}\end{aligned}\tag{19}$$

$$\begin{aligned}i_t &= [i - (y_c - r)] * s_y + t_y \\j_t &= [j - (x_c - r)] * s_x + t_x\end{aligned}$$

The relationship in (19) provides a method for putting any part of an image into the circle at a reasonable scale.

Now what if we let the user arbitrarily pick four points to define a mapping from the texture images to the circle? It becomes much more difficult to map a square onto an arbitrary set of four unique points if we try to do it geometrically. However, we can cheat.

2.7.4 Interpolation

Given a circle, its bounding box is a square in the image.

- Pick four points that define a quadrilateral in the texture image
- Assign the four points to the four points of the square
- To discover the texture coordinate of a pixel in the circle, use bilinear interpolation.

Interpolation is a method of smoothly varying a value from one state to another. Linear interpolation makes the rate of change constant. Bilinear interpolation is useful on 2D surfaces like images, because it makes use of four endpoints instead of two.

The process for bilinear interpolation is as follows.

- Given four corners A, B, C, D and a location (i, j) in between them
- Based on the column j , mix the values of corners A and B , put the result in E
- Based on the column j , mix the values of corners C and D , put the result in F
- Based on the row i , mix the values of E and F

The end result is a texture coordinate for the pixel (i, j) .

What if the pixel within the circle covers many pixels in the texture image?

- Identify the texture coordinates of the four corners of the pixel: $(i, j), (i+1, j), (i, j+1), (i+1, j+1)$
- Calculate the enclosing bounding box of the four texture coordinates
- Average the pixel values within the bounding box

2.7.5 Scanline Fill Algorithm

The scanline fill algorithm is one of the most commonly used algorithms in computer graphics. Its use goes far beyond simply filling polygons with a single color. The basic concept is the underlying algorithm for many rendering systems.

The basic concept is simple:

- For each scanline, intersect the scanline with all the edges of the polygon
- Sort the edges left to right and fill between the edge pairs

Problems can arise, however, because some vertices require special handling

- Some vertices need to be counted twice
- Some vertices need to be counted once

It turns out the the property differentiating the two cases is whether the vertex is an extremum or not. Vertices that are also vertical extrema need to be counted twice. Vertices that are not extrema should only be counted once. Note that, in the case of non-extrema vertices, we could also just shorten one of the edges by one pixel.


Rather than calculating the intersections of the polygon edges with the scanlines explicitly, we can take advantage of **coherence** since the edges of the polygon are lines. The slopes don't change. Therefore, we can just keep track of the current boundary point for each edge that intersects the current scanline and update those in fixed increments as we move to the next scanline. Then the major issue is just handling which edges intersect the current scanline.

We could use Bresenham's algorithm to calculate the boundary pixels, but it turns out there is no advantage to not using floating point math for polygons: most of the time we're trying to render polygons that are mappings of 3D objects onto a 3D plane and we want to be as accurate as possibly with our placement.

Ordered Edge List Algorithm

1. Process a polygon by going through the vertices in order
2. For each non-horizontal edge
 - Calculate edge properties (e.g. inverse slope, start point, end point)
 - Insert the edge into an ordered edge list according to starting scanline
 - Set the endpoint properly, shortening it by one if it is not an extremum endpoint.
 - Keep track of the starting scanline
3. For each scanline until the edge list is empty
 - Put any edges that start on this scanline into the active edge list (keep it sorted)
 - Fill the scanline from left to right, pairing up edges
 - Remove edges that terminate on this scanline
 - Update the active edges by incrementing their intersections
 - Sort the active edge list using the new intersection values

Polygon Example



Edges:

- E_1 : (0, 0) (4, 0)
- E_2 : (4, 0) (4, 3) [modified to (4, 0) (4, 2) in preprocessing step]
- E_3 : (4, 3) (0, 5)
- E_4 : (0, 5) (0, 0)

Ordered edge list: HEAD $\rightarrow E_4(0, 4) \rightarrow E_2(0, 2) \rightarrow E_3(3, 4) \rightarrow \text{NULL}$

Active Edge List:

- Scanlines {0, 1, 2}: HEAD $\rightarrow E_4(0, 4) \rightarrow E_2(0, 2) \rightarrow \text{NULL}$
- Scanlines {3, 4}: HEAD $\rightarrow E_4(0, 4) \rightarrow E_3(3, 4) \rightarrow \text{NULL}$

Going through the process from scanline 0 to 5 draws the shape above. Shortening E_2 solves the double-counting problem, but the polygon is way too big.

2.7.6 Matching polygons to the grid

The following modifications repair the algorithm so it creates polygons with the correct area, if they are integer coordinates.

- Shorten the top edge of all edges by 1 (not just non-extrema)
- Fill up to the right pixel, but don't color the last one
- Add $\frac{1}{2} \frac{dx}{dy}$ to the initial intersection

For non-integer coordinates, it is important to account for the exact location of the vertex within the pixel. The general rule is that the x-intersection needs to be initialized to the intersection of the edge with the first scanline below it. If V_y is the smaller Y coordinate, the following cases determine the modifier to the initial x-intersect.

- Case $V_y - \text{floor}(V_y) \leq 0.5$: add $(0.5 - (\text{floor}(V_y) - V_y)) \frac{dx}{dy}$
- Case $V_y - \text{floor}(V_y) > 0.5$: add $((1.0 - (\text{floor}(V_y) - V_y)) + 0.5) \frac{dx}{dy}$

2.8 Interpolation and Barycentric Coordinates

The scanline fill algorithm provides a useful framework for not only drawing a polygon, but also shading it, texturing it, or otherwise modifying it based on the properties of its vertices. If, for example, we know the color of the polygon at each vertex, then we can interpolate the color along each edge and then interpolate the color along each scanline.

Interpolation is the process of smoothly blending from one value to another across some distance. In this case, the distance is across pixels. In general, if you have a value C_A at location x_0 and a value C_B at location x_1 , then you can use a variation of the alpha blending equation to implement linear interpolation from C_A to C_B across the distance from x_0 to x_1 . At each location x , the new color $C(x)$ is defined as in (20).

$$C(x) = C_A \left(1 - \frac{x - x_0}{x_1 - x_0} \right) + C_B \left(\frac{x - x_0}{x_1 - x_0} \right) \quad (20)$$

If we are using linear interpolation, then the change in C for a fixed size step in x is constant.

It is also possible to do interpolation across a plane as a weighted combination of three vertices A , B , and C . The weights for the vertices A, B, C are given by the barycentric coordinates (α, β, γ) , which we can compute for any (x, y) . The barycentric coordinates must obey the relationships in (21), where $P_{x,y}$ is some point on the plane.

$$\begin{aligned} \alpha A + \beta B + \gamma C &= P_{x,y} \\ \alpha + \beta + \gamma &= 1 \end{aligned} \quad (21)$$

Conceptually, barycentric coordinates are defined by a possibly non-orthogonal coordinate system. In practice, we tend to use three points—the vertices of a triangle—and are mostly concerned about the coordinates of pixels inside the triangle. Given three vertices, (A, B, C) , the barycentric coordinates are defined by the perpendicular distance from the line opposite the vertex. The distance from line CA is the coordinate for point B ; the distance from line AB is the coordinate for point C , and the distance from line BC is the coordinate for point A . Note the orientation of the line segments: the traversal of the vertices is in

We can use the implicit algebraic equation for a line to calculate the barycentric coordinates. The implicit equation for a line generated by two points $(x_0, y_0), (x_1, y_1)$, is given by (22).

$$f(x, y) = (y_0 - y_1)x + (x_1 - x_0)y + x_0y_1 - x_1y_0 \quad (22)$$

Using the implicit equation, we can calculate β and γ and then use (21) to compute α .

$$\begin{aligned} \beta &= \frac{f_{ca}(x, y)}{f_{ca}(x_b, y_b)} = \frac{(y_c - y_a)x + (x_a - x_c)y + x_cy_a - x_ay_c}{(y_c - y_a)x_b + (x_a - x_c)y_b + x_cy_a - x_ay_c} \\ \gamma &= \frac{f_{ab}(x, y)}{f_{ab}(x_c, y_c)} = \frac{(y_a - y_b)x + (x_b - x_a)y + x_ay_b - x_by_a}{(y_a - y_b)x_c + (x_b - x_a)y_c + x_ay_b - x_by_a} \\ \alpha &= 1 - \beta - \gamma \end{aligned} \quad (23)$$

2.9 Anti-aliasing

Aliasing is caused by sampling a continuous world at too low a frequency. Aliasing can occur in many different forms.

- Spatial aliasing: trying to represent continuous objects using a discrete grid
- Time aliasing: trying to represent continuous motion using a discrete time step
- Spectral aliasing: trying to represent continuous spectra using discrete color bands

A good graphics system must deal with all three (the third we don't tend to notice so much because the sampling frequency is high enough for most situations).

Spatial aliasing is most relevant for object primitives. There are a number of methods for improving the appearance of objects.

2.9.1 Supersampling

The basic concept of supersampling is to render the object onto a much larger temporary image and use the higher resolution rendering to calculate how to draw the object into the original image.

- Use 9 or 16 subpixels per final pixel
- Draw the object scaled up appropriately
- Count how many subpixels are on within each final pixel to draw into the final image
- Alpha blend when putting the object into the final image

You can also draw the entire scene in an image that is large and then scale it down. When working with polygons, this is fine and works well. For lines, however, unless you draw the lines several pixels thick in the high resolution image, they will almost disappear in the final image.

One problem with supersampling single lines is that a line can intersect at most 3 subpixels within a superpixel. Therefore, you have only four levels of coloring $[0, \frac{1}{3}, \frac{2}{3}, 1]$.

2.9.2 Area Sampling

The concept of area sampling is to use an expanded canvas, but draw the line with appropriate thickness. If you're using a 3x3 grid, for example, then draw a 3-pixel wide line.

- Keep track of the upper and lower edges of the line using Bresenham's algorithm (left and right for 2nd and 3rd octants)
- Fill between the lines
- Add a little bit for the ends

The distance between the upper and lower lines depends upon the slope, as shown in figure 4.

$$y_{\text{upper}} = mx + b + \frac{1}{2}\sqrt{m^2 + 1}$$

$$y_{\text{lower}} = mx + b - \frac{1}{2}\sqrt{m^2 + 1}$$
(24)

The benefit of using thick lines and then supersampling is that the lines don't disappear. Because all of the subpixels in a pixel could be filled, you get the full benefit of increasing the size of the image. When you supersample at 3x3, you get 10 levels of brightness.




Figure 4: A solid line in an upsampled image

2.9.3 Error Terms for Antialiasing

How about using the error term in Bresenham's algorithm to antialias a line?

- The error term is $2dx$ times the difference in the distances between pixel centers
- If the error term is in the range $[0, 1.0]$ then the fraction to color the top pixel is $\alpha = 1 - \frac{e}{2}$
- If the error term is in the range $[-1.0, 0)$ then color the top pixel $\alpha = \frac{1}{2} + \frac{e}{2}$
- If the error term is in the range $[-2.0, -1.0)$ then color the pixel below $\alpha = -\frac{e}{2} - \frac{1}{2}$
- The fraction for the current pixel is always $f = 1 - \alpha$.

Since the scanline-fill algorithm maintains the x value of the intersection with the current scanline and the slope, you have all the information you need to properly color the current pixel of a polygon—use the fraction of the pixel covered by the polygon. Note that, to properly anti-alias, you must make the calculation for each pixel, and only pixels completely below the starting or ending edge should be fully lit.

When rendering a scene with a polygonal mesh, you probably don't want to anti-alias individual polygons. Supersample instead.

2.9.4 Post-filtering

If you draw your scene in an image that is much larger than the final one (using appropriate width lines), then you can apply a post-filtering process to reduce the size of the image. A post-filter is generally a form of **convolution**. Convolution is simply the set of rules for how to take an image A and filter it with another image B. Generally, B is a small image that represents some operation we want to execute on A.

The algorithm for convolution of image A by image B is as follows.

Create an output image Z the size of A.

For each pixel (i, j) in A

Place B so its center is over pixel (i, j)

Initialize a sum variable S to 0

For each pixel (x, y) in B

Multiply pixel (x, y) and its corresponding pixel in A

Add the product to S

Assign S to pixel (i, j) of Z

It is important to note that, technically, B should be flipped horizontally and vertically before running the algorithm above. For any filter with horizontal and vertical symmetry, however, it makes no difference.

Common filters for anti-aliasing are box filters and Gaussian filters.

- Box filter - the values of the filter image B are all the same and sum to 1 (simple averaging)
- Gaussian filter - the values of the filter image B are calculated using a 2D Gaussian distribution (normal distribution)

Real applications (including operating systems) draw anti-aliased lines. The jaggies are bad.

2.10 Cells, Icons, Sprites, and Characters

Cell arrays are small 2-D arrays of pixels. In addition to colors, they can hold an alpha map.

- Icons: generally square 8x8, 16x16, or 32x32
- Sprites: arbitrarily scaled rectangles, generally with an alpha map, often used for 2D characters in games
- Characters: fonts used to be represented by cell arrays containing a binary mask

Most graphics systems give you lots of options for plotting cell arrays on a screen.

- Source copy: paste it over the current image content
- Source mix: use a single alpha value (in addition to any inherent alpha map)
- Source invert: use the map to invert the current image content
- Logical operations: can do bitwise AND, OR, NOT, and XOR operations

3 2-D Transformations and Viewing

When describing a scene or animation, we want to be able to describe transformations to an object that cause it to appear a certain way. Common transformations include translation, rotation, and scaling.

If we want to translate an object around the screen, we have three options.

1. Change the coordinates by hand
2. Add an offset to the current coordinates
$$\begin{bmatrix} x_n \\ y_n \end{bmatrix} = \begin{bmatrix} x_0 \\ y_0 \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$
3. Develop a general system of representing and applying transformations

If we want to rotate an object around the screen, we have three options.

1. Change the coordinates by hand
2. Transform the points along the path of a circle
$$\begin{bmatrix} x_n \\ y_n \end{bmatrix} = \begin{bmatrix} x_0 \cos \theta - y_0 \sin \theta \\ x_0 \sin \theta + y_0 \cos \theta \end{bmatrix}$$
3. Develop a general system of representing and applying transformations

If we want to scale an object, we have three options.

1. Change the coordinates by hand
2. Multiply each point by a scalar value
$$\begin{bmatrix} x_n \\ y_n \end{bmatrix} = \begin{bmatrix} x_0 s_x \\ y_0 s_y \end{bmatrix}$$
3. Develop a general system of representing and applying transformations

While option B works fine, option C is much easier and permits easy manipulation and mixing of transformations.

3.1 Matrix Representations

Homogeneous coordinates are the basis for building a generic manipulation system. The homogeneous coordinates for a 2-D Cartesian point (x_c, y_c) are (x, y, h) . The relationship between the homogeneous and standard coordinates is given in 25.

$$\begin{aligned} x_c &= \frac{x}{h} \\ y_c &= \frac{y}{h} \end{aligned} \tag{25}$$

Note that this permits many representations of a single 2-D Cartesian point. In the standard, or normalized form of homogeneous coordinates $h = 1$. Using homogeneous coordinates lets us implement transformations such as translation, rotation, and scaling using a single matrix representation.

Translation

$$\begin{bmatrix} x_n \\ y_n \\ 1 \end{bmatrix} = \begin{bmatrix} x_0 + t_x \\ y_0 + t_y \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_0 \\ y_0 \\ 1 \end{bmatrix} \tag{26}$$

Rotation (about Z-axis)

$$\begin{bmatrix} x_n \\ y_n \\ 1 \end{bmatrix} = \begin{bmatrix} x_0 \cos \theta - y_0 \sin \theta \\ x_0 \sin \theta + y_0 \cos \theta \\ 1 \end{bmatrix} = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_0 \\ y_0 \\ 1 \end{bmatrix} \quad (27)$$

Scaling

$$\begin{bmatrix} x_n \\ y_n \\ 1 \end{bmatrix} = \begin{bmatrix} x_0 s_x \\ y_0 s_y \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_0 \\ y_0 \\ 1 \end{bmatrix} \quad (28)$$

Because each 2-D transformation is represented by a 3x3 matrix, we can use matrix multiplication and manipulation to concatenate transformations. Since we use column vectors for our coordinates, the order of operation moves right to left. The right-most transformation executes first.

- Order doesn't matter when concatenating only translations, only rotations, or only scales.
- Order matters when you mix different transformation type.

By concatenating a series of transformations, we can do things like rotate or scale an object around its center.

Rotating around a pivot

1. Translate from the pivot point (x_p, y_p) to the origin
2. Rotate by θ around the origin
3. Translate back to the pivot point.

$$\begin{aligned} R_{x_p, y_p}(\theta) &= T(x_p, y_p) R_z(\theta) T(-x_p, -y_p) \\ &= \begin{bmatrix} 1 & 0 & x_p \\ 0 & 1 & y_p \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & -x_p \\ 0 & 1 & -y_p \\ 0 & 0 & 1 \end{bmatrix} \\ &= \begin{bmatrix} \cos \theta & -\sin \theta & x_p(1 - \cos \theta) + y_p \sin \theta \\ \sin \theta & \cos \theta & y_p(1 - \cos \theta) + x_p \sin \theta \\ 0 & 0 & 1 \end{bmatrix} \end{aligned} \quad (29)$$

Scaling around a pivot

1. Translate from the pivot point (x_p, y_p) to the origin
2. Scale around the origin by (s_x, s_y)
3. Translate back to the pivot point

$$\begin{aligned}
R_{x_p, y_p}(\theta) &= T(x_p, y_p)R_z(\theta)T(-x_p, -y_p) \\
&= \begin{bmatrix} 1 & 0 & x_p \\ 0 & 1 & y_p \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & -x_p \\ 0 & 1 & -y_p \\ 0 & 0 & 1 \end{bmatrix} \\
&= \begin{bmatrix} s_x & 0 & x_p(1 - s_x) \\ 0 & s_y & y_p(1 - s_y) \\ 0 & 0 & 1 \end{bmatrix}
\end{aligned} \tag{30}$$

Scaling around a pivot and orientation

1. Translate from the pivot point (x_p, y_p) to the origin
2. Rotate to align the scale axis $\hat{n} = (n_x, n_y)$ with the x-axis $R(-\theta) : \theta = \arctan \frac{n_y}{n_x}$
3. Scale about the origin by (s_x, s_y)
4. Rotate back by θ
5. Translate back to the pivot point

$$S_{(x_p, y_p, \hat{n})}(s_x, s_y) = T(x_p, y_p)R(\theta)S(x_p, y_p)R(-\theta)T(-x_p, -y_p) \tag{31}$$

You can pre-compute the complete matrix or just let the computer multiply the matrices.

Shears

A shear moves different parts of the coordinate system differently. Shears can simulate tilting of the 2-D coordinate plane away from the viewer.

$$Sh_x(k_x) = \begin{bmatrix} 1 & k_x & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \tag{32}$$

$$Sh_y(k_y) = \begin{bmatrix} 1 & 0 & 0 \\ k_y & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \tag{33}$$

$$Sh_{xy}(k_x, k_y) = \begin{bmatrix} 1 & k_x & 0 \\ k_y & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \tag{34}$$

3.2 2-D Viewing Pipeline

Consider the following situation. You have designed a nice little 2D logo. You want to make the logo appear as though it were on the label of a vinyl record. Furthermore, the record is on a turntable that is spinning. You want the camera view to begin with the turntable fully visible, then zoom in on the label, moving with it as the view gets closer until the label appears to be rotating in place in the middle of the image as the rest of the scene rotates around it. All of this should be rendered on a 4500px by 8000px image.

For any given frame of the sequence, the transformation from model point to screen coordinate is describable by a single matrix transformation. However, building the transformation is a non-trivial process, and there are many parameters. We want to control the location of the label on the record, the speed of the turntable, what part and what amount of the world is visible, and the size of our final image.

Breaking apart the scenario in different coordinate systems is an extremely useful way to subdivide the problem into manageable pieces. In fact, it is so useful, that there is a standard pipeline of coordinate systems that almost all CG scenes use to move from model coordinates to screen coordinates.

1. **Model Coordinates** - Model coordinates are where you do your work. The model coordinate frame is attached to the object in a convenient way.
 - Each model can have its own coordinate system
 - A collection of models can have its own coordinate system
 - Model coordinates should make it easy to describe the model or collection

In the example above, there are several useful model coordinate systems.

- The logo should have its own origin and orientation that make it easy to describe. The central rotation point of the logo, for example, would be a useful model origin.
 - The record should have its own origin. The center of the record makes the most sense, in this case. The orientation of the axes are arbitrary since the object is radially symmetric.
 - The turntable should have its own model coordinates; a corner origin makes it simple to define.
2. **World Coordinates** - The world coordinate frame allows you to place models in relation to one another, providing a single frame of reference. The world coordinate system is the frame in which the viewing system is defined.
 - Pick a convenient origin
 - A transformation defines the location of each instance of each model in the scene
 - We can create multiple instances of each model or object
 - Creating a new transformation matrix for an object puts it in a different place
 - Duplicating the polygons in a model and adjusting the transformation makes a new copy

In the example above, aligning the world coordinates with the turntable coordinates is not a bad choice. It is the largest object in the scene, and locating it at (0,0) makes sense. If the turntable were itself sitting on a table inside a house, then choosing a world coordinate system aligned with the house would be more appropriate. Given the world coordinate system, we can specify the location of each object. Note that, in some situations, it's better to define the position of objects relative to one another—e.g. the record relative to the turntable—rather than using world coordinates.

3. Viewing Coordinates - The viewing coordinate frame specifies how the camera views the world. The view coordinates are defined relative to world coordinates through an imaging model.

- In 2D, the imaging model is generally a box, but other shapes are possible
- The box can be located anywhere, possibly in any orientation
- Need to specify the size, location, and orientation of the box
- Viewing coordinates are defined relative to world coordinates so you know where the camera is located relative to objects in the scene.

In the example above, the 2D view window could be defined by an origin in the center of the window, the length of its sides (in world coordinate units) and the orientation of the view window's horizontal axis relative to the world coordinate x-axis. It is important to think about the view window in world coordinate units, not pixels. The number of pixels in the final image is arbitrary and many different size images could be generated from a single view of a scene.

To move from world coordinates to viewing coordinates, we need to translate the origin of the view window (v_x, v_y) to the world coordinate origin and then rotate to orient the view window x-axis (n_x, n_y) with the world coordinate x-axis. If (n_x, n_y) represent a normalized vector, then the transformation is given below.

$$X_{\text{world} \rightarrow \text{view}}(v_x, v_y, n_x, n_y) = R_z(n_x, -n_y)T(-v_x, -v_y) \quad (35)$$

In general, if you have a coordinate system defined by two orthonormal vectors (u_x, u_y) and (v_x, v_y) , then the matrix to align the coordinate system with the world axes is given below.

$$X_{\text{align}}(\vec{u}, \vec{v}) = \begin{bmatrix} u_x & u_y & 0 & 0 \\ v_x & v_y & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad (36)$$

4. Normalized Viewing Coordinates - Normalized viewing coordinates are independent of the input or output viewing system. The points are scaled to be in the range $[0, 1]$
- Clipping is possible (and fast) with normalized viewing coordinates
 - Normalized coordinates are ready for multiple display systems

To move from viewing coordinates to normalized viewing coordinates we scale by the width du and height dv of the view window.

$$X_{\text{view} \rightarrow \text{norm}}(du, dv) = S\left(\frac{1}{du}, \frac{1}{dv}\right) \quad (37)$$

5. Screen Coordinates - Screen coordinates specify which pixels to light up. The coordinate system is defined by the output display device and each output display device will have its own. Using normalized viewing coordinates makes it easy to send the output to numerous different output devices.

To move to screen coordinates we need to invert the y-axis, scale to the number of pixels in each direction, then translate the view window so all y-coordinates are positive.

$$X_{\text{norm} \rightarrow \text{screen}}(\text{rows}, \text{cols}) = T(\text{cols}/2, \text{rows}/2)S(\text{cols}, -\text{rows}) \quad (38)$$

4 3-D Transformations

Move from 2-D to 3-D adds a z-coordinate to the points and matrices. Otherwise, the concepts are identical. The number of rotations increases, as we can rotate around any of the three axes, not just the z-axis.

Translation

$$\begin{bmatrix} x_n \\ y_n \\ z_n \\ 1 \end{bmatrix} = \begin{bmatrix} x_0 + t_x \\ y_0 + t_y \\ z_0 + t_z \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & t_x \\ 0 & 1 & 0 & t_y \\ 0 & 0 & 1 & t_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_0 \\ y_0 \\ z_0 \\ 1 \end{bmatrix} \quad (39)$$

Rotation (about X-axis)

$$\begin{bmatrix} x_n \\ y_n \\ z_n \\ 1 \end{bmatrix} = \begin{bmatrix} x_0 \\ y_0 \cos \theta - z_0 \sin \theta \\ y_0 \sin \theta + z_0 \cos \theta \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & \cos \theta & -\sin \theta & 0 \\ 0 & \sin \theta & \cos \theta & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_0 \\ y_0 \\ z_0 \\ 1 \end{bmatrix} \quad (40)$$

Rotation (about Y-axis)

$$\begin{bmatrix} x_n \\ y_n \\ z_n \\ 1 \end{bmatrix} = \begin{bmatrix} x_0 \cos \theta + z_0 \sin \theta \\ y_0 \\ -x_0 \sin \theta + z_0 \cos \theta \\ 1 \end{bmatrix} = \begin{bmatrix} \cos \theta & 0 & \sin \theta & 0 \\ 0 & 1 & 0 & 0 \\ -\sin \theta & 0 & \cos \theta & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_0 \\ y_0 \\ z_0 \\ 1 \end{bmatrix} \quad (41)$$

Rotation (about Z-axis)

$$\begin{bmatrix} x_n \\ y_n \\ z_n \\ 1 \end{bmatrix} = \begin{bmatrix} x_0 \cos \theta - y_0 \sin \theta \\ x_0 \sin \theta + y_0 \cos \theta \\ z_0 \\ 1 \end{bmatrix} = \begin{bmatrix} \cos \theta & -\sin \theta & 0 & 0 \\ \sin \theta & \cos \theta & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_0 \\ y_0 \\ z_0 \\ 1 \end{bmatrix} \quad (42)$$

Scaling

$$\begin{bmatrix} x_n \\ y_n \\ z_n \\ 1 \end{bmatrix} = \begin{bmatrix} x_0 s_x \\ y_0 s_y \\ z_0 s_z \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 & 0 \\ 0 & s_y & 0 & 0 \\ 0 & 0 & s_z & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_0 \\ y_0 \\ z_0 \\ 1 \end{bmatrix} \quad (43)$$

Rotations about a single axis commute, but not rotations about different axes. One other interesting fact about rotations is that to reverse the effect of a rotation you just transpose the matrix.

$$R_a(-\theta) = R_a^t(\theta) \quad (44)$$

The transpose of a rotation matrix is also its inverse, which makes sense, as $R_a(\theta)R_a(-\theta) = I$.

Orienting Coordinate Systems

If you have an orthonormal basis (three orthogonal, normalized vectors that follow the right hand rule), then rotating the basis to orient with the primary axes is a one-step process. Given the three basis vectors $\vec{u}, \vec{v}, \vec{w}$, the matrix that orients the coordinate system with the primary axes is given by $R_{xyz}(\vec{u}, \vec{v}, \vec{w})$. The transpose of the matrix will execute the inverse rotation.

$$R_{xyz}(\vec{u}, \vec{v}, \vec{w}) = \begin{bmatrix} u_x & u_y & u_z & 0 \\ v_x & v_y & v_z & 0 \\ w_x & w_y & w_z & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad (45)$$

The orienting matrix makes it possible to easily do things like scale in an arbitrary coordinate system ($\vec{u}, \vec{v}, \vec{w}$) around a pivot point (p_x, p_y, p_z) .

$$S(\vec{u}, \vec{v}, \vec{w}; p_x, p_y, p_z; s_x, s_y, s_z) = T(p_x, p_y, p_z) R^t(\vec{u}, \vec{v}, \vec{w}) S(s_x, s_y, s_z) R(\vec{u}, \vec{v}, \vec{w}) T(-p_x, -p_y, -p_z) \quad (46)$$

Shears, mirroring, and other transformations are also possible in 3D, just as in 2D. More complex shears are possible because any two axes can affect the third.

5 Viewing

Generic 2D viewing requires only rotations, translations, and scales in order to implement a view rectangle at an arbitrary location and orientation. Since we are converting coordinates from a plane to a plane, we don't need anything special in order to view the scene.

In order to view a 3D scene, however, we have to project 3D points onto a plane. Mappings onto a plane are the domain of plane geometric projections. Generic geometric projects are, in general, between equivalent dimensional spaces. 3D rotations, translations, and scales map from one 3D space to another. Plane geometric projects, however, drop one dimension. Therefore, the rank of a 4x4 plane geometric projection matrix can be no more than three.

Geometric projections fall into one of two categories.

- Affine: affine transforms maintain parallel lines
- Perspective: perspective transforms do not necessarily maintain parallel lines (e.g. railroad tracks going off into the distance)

When defining viewing pipelines, we can view the process in one of two ways.

- Fixed viewing geometry: the viewer remains stationary and the object moves
- Fixed object geometry: the scene remains stationary and the viewer moves

Both approaches produce identical matrices, but the two processes make sense in different situations. A fixed viewing geometry is useful when working with CAD models or 3D modeling programs. Then the situation is similar to a carver holding a statue and moving it around in front of their eyes. Fixed object

geometry is useful when rendering a scene from the point of view of someone within it. We don't tend to think of a city flowing by us or rotating around as we walk down a sidewalk.

Planar geometric projections divide into two large groups and then several subcategories within them. Most of the subdivisions are specialty viewing situations designed to be useful in engineering drawings.

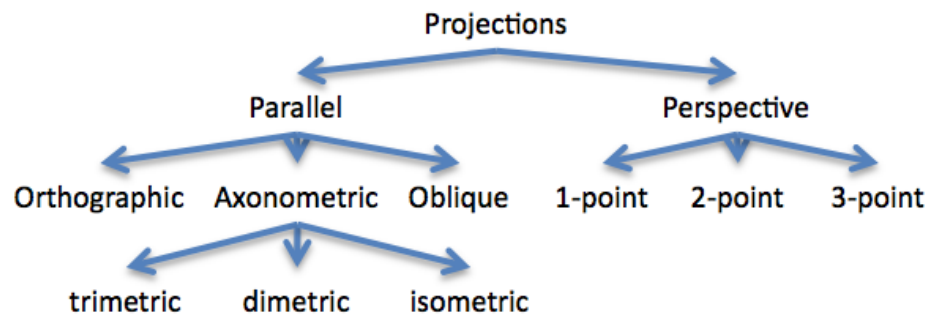


Figure 5: Categories of planar projections

All projection methods can be defined by a plane and a center of projects [COP]. To find the location of a 3D world point on the view plane, create a line from the center of projection through the world point and find where it intersects with the view plane. That defines its projection. The difference between parallel and perspective projections is whether the COP is at infinity (parallel) or at a finite distance (perspective).

5.1 Orthographic Projection

Orthographic projection is the simplest of the projection methods: drop the z-coordinate. Conceptually, the view plane is a rectangle parallel to the x-y plane. The COP is at infinity in the Z-direction, so each line perpendicular to the x-y plane is a line of projection. The orthographic projection matrix is an identity matrix with the z-row set to all zeros.

$$P_{\text{orth}} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad (47)$$

Engineering drawings are normally created using orthographic projection. By using the six possible combinations of axes and COP orientations an object can be rendered so all faces are visible in at least one drawing. Furthermore, distances are preserved for any face parallel to the projection plane.

5.2 Axonometric Projections

Axonometric projection is a variation of orthographic projection. The various axonometric projections occur when multiple faces of an object are visible. The concept is only meaningful for fixed viewing geometry. The subdivisions of axonometric viewing indicate how many rulers are required to measure the different faces of the object.

1. Trimetric projection: three faces of an object are visible, but the three faces require different rulers. When a line segment is angled relative to the view plane, its projected length is dependent upon the angle. A larger angle means the same unit step in the image corresponds to a longer step in the 3D scene. In trimetric projection, the three sides of the object have unique angles relative to the view plane.
2. Dimetric projection: two faces of an object share the same ruler. Two of the a sides of the object have identical angles relative to the view plane.
3. Isometric projection: all faces of the object share the same ruler. There are only four possible isometric projections.
 - The foreshortening factor is 0.8165 (a step of 0.8165 in the world corresponds to a step of 1 in the image).
 - The angle with the horizontal axes and the projected axis is $\pm 30^\circ$.
 - Isometric projections are often used in engineering drawings because it is easy to measure distances.

5.3 Oblique Projections

Orthographic and axonometric projections both use lines of projection that are perpendicular to the view plane. Oblique projections occur when the line from the view plane origin to the COP is not perpendicular to the view plane. Conceptually, it's like staring straight ahead and seeing stuff to the side (but not stuff straight ahead). With oblique projections, faces parallel to the view plane are rendered without distortion, but non-parallel faces and lines are distorted (basically with a shear).

Oblique projection lets you see the side of an object when that side is perpendicular to the view plane.

Cavalier projection: the angle between the projective lines and the view plane is $\alpha = 45^\circ$.

- The foreshortening factor for lines perpendicular to the view plane is $\sqrt{2} = 0.7071$

Cabinet projection: the angle between the projective lines and the view plane is $\alpha = 63.43^\circ$.

- The foreshortening factor for lines perpendicular to the view plane is 0.5

In general, f is the desired foreshortening factor for the z-axis and α is the angle the projected z-axis makes with the x-axis. The horizontal orientation of the projection vector is determined by f . The vertical orientation of the projection vector is determined by α .

$$P_{\text{oblq}} = \begin{bmatrix} 1 & 0 & -f \cos \alpha & 0 \\ 0 & 1 & -f \sin \alpha & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad (48)$$

5.4 Arbitrary Parallel Projections

All of the above projections are suitable for a fixed viewing geometry situation. However, we can also set up parallel projects for fixed scene geometry in a way that includes all of the above projection methods. The basic concept is to define a camera in the world and then develop a viewing pipeline that transforms world coordinates into image coordinates.

The camera:

- View Reference Point [VRP] is the origin (center) of the view plane $VRP = (P_x, P_y, P_z)$
- View Plane Normal [VPN] is the normal to the view plane $VPN = (n_x, n_y, n_z)$
- View Up Vector [VUP] is the direction of the y-axis on the view plane $VUP = (Y_x, Y_y, Y_z)$
- View plane extent in coordinates relative to the VRP on the view plane $(u_0, v_0), (u_1, v_1)$
- Direction of Projection [DOP] is the direction of the lines of projection in the coordinate system defined by the VPN and VUP vectors $DOP = (D_x, D_y, D_z)$
- Front clip plane and back clip plane, specified as distances along the VPN, F , and B .
- Screen size in pixels R and C .




Figure 6: View geometry showing the view plane, its origin (VPN), the view reference coordinates (VPN, VUP, and \hat{U}), the direction of project (DOP) and the view window boundaries $(U_{\min}, U_{\max}), (V_{\min}, V_{\max})$

The basic idea of the parallel projection process is to align the view reference coordinate system with the world coordinate axes, shear so the view volume is aligned with the axes, scale the volume to an appropriate size, and then project the world onto a plane using orthographic projection.

The process:

1. Calculate the U vector for the view reference coordinate system

$$U = \text{VUP} \times \text{VPN} \quad (49)$$

2. Recalculate the VUP so that the view reference coordinates are orthogonal

$$\text{VUP} = \text{VPN} \times U \quad (50)$$

3. Translate the VRP to the origin

$$\text{VTM}_1 = T(-\text{VRP}) \quad (51)$$

4. Normalize U , VUP and VPN and use the R_{xyz} rotation method to align the axes.

$$\text{VTM}_2 = R_{xyz}(\hat{U}, \hat{\text{VUP}}, \hat{\text{VPN}})\text{VTM}_1 \quad (52)$$

5. Shear so the DOP is aligned with the positive z-axis.

$$\text{VTM}_3 = Sh_z\left(\frac{D_x}{D_z}, \frac{D_y}{D_z}\right)\text{VTM}_2 \quad (53)$$

$$Sh_z(k_x, k_y) = \begin{bmatrix} 1 & 0 & k_x & 0 \\ 0 & 1 & k_y & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad (54)$$

6. Translate so that the center of the view window and the front clip plane are at the origin.

$$\text{VTM}_4 = T\left(-\frac{u_0 + u_1}{2}, -\frac{v_0 + v_1}{2}, -F\right)\text{VTM}_3 \quad (55)$$

7. Scale to the canonical view volume, which is the box bounding by $[-1, 1]$ in x and y and $[0, 1]$ in z .

- The distance between the front and back clip planes is $B' = B - F$
- The width of the view window is $du = u_1 - u_0$
- The height of the view window is $dv = v_1 - v_0$

$$\text{VTM}_{\text{World} \rightarrow \text{CVV}} = S\left(\frac{2}{du}, \frac{2}{dv}, \frac{1}{B'}\right)T\left(-\frac{u_0 + u_1}{2}, -\frac{v_0 + v_1}{2}, -F\right)Sh_z\left(\frac{D_x}{D_z}, \frac{D_y}{D_z}\right)R_{xyz}(\hat{U}, \hat{\text{VUP}}, \hat{\text{VPN}})T(-\text{VRP}) \quad (56)$$

At this point in the process, all vertices that are visible are within the canonical view volume. Clipping is normally executed here, since tests within the bounding volume are simple to calculate and evaluate. Clipping to the front and back of the view volume are especially important, as they cannot be handled in a 2D clipping process.

8. The next step in the process is to use orthographic projection to convert the 3D points to 2D points on the plane. Since the canonical view volume is aligned with the major axes and the DOP is aligned with the z-axis, orthographic projection just drops the z-coordinate.
9. The final step in the process is to scale and translate the scene into screen coordinates. Note that because we used a right-handed coordinate system for the view reference coordinates and made positive Z our viewing axis, that the X-axis points to the left. So we need to invert both the x- and y-axes to get screen coordinates.

$$\text{VTM}_{\text{CVV} \rightarrow \text{Screen}} = T\left(\frac{C}{2}, \frac{R}{2}\right)S\left(-\frac{C}{2}, -\frac{R}{2}\right)P_{\text{orth}} \quad (57)$$

5.5 Perspective Projection

Parallel projection is often used in 3D modeling systems (e.g. CAD Design) and some video games because of the need to map from 2D back into 3D (mouse clicks to object creation) and because of the need to compare relative distances accurately. Our eyes, however, do not implement parallel projection. Instead, things that are farther away get smaller.

Perspective projection was one of the concepts that came out of the Renaissance. Artists often think about perspective viewing in a fixed-view sense, orientating objects in the scene according to a 1-point, 2-point, or 3-point perspective view.

- A 1-point perspective occurs when the object's front face is parallel to the view plane, which means parallel lines on the front face remain parallel.
- A 2-point perspective occurs when the object is oriented so that lines in only one axis (usually the y-axis) are parallel to the view plane. For example, when a cube is oriented so an edge is towards the viewer and the line along the edge is parallel to the view plane.
- 3-point perspective occurs when none of the major object edges are parallel to the view plane.

In graphics, however, the fixed scene geometry view makes more sense with perspective projection. A useful paradigm is to think of the viewer as a camera that can move around and through the scene. Similar to a camera, we can parameterize the view with values that correspond to the size of the lens and the zoom setting.

The camera:

- View Reference Point [VRP] is the origin (center) of the view plane $\text{VRP} = (P_x, P_y, P_z)$
- View Plane Normal [VPN] is the normal to the view plane $\text{VPN} = (n_x, n_y, n_z)$
- View Up Vector [VUP] is the direction of the y-axis on the view plane $\text{VUP} = (Y_x, Y_y, Y_z)$
- View plane extent in coordinates relative to the VRP on the view plane $(u_0, v_0), (u_1, v_1)$
- Center of Projection [COP] is the location of the focal point of the projection, given in View Reference Coordinates $(\text{COP}_x, \text{COP}_y, \text{COP}_z)$. In many cases, the COP is reduced to an offset d along the negative VPN axis.
- Front clip plane and back clip plane, specified as distances along the VPN, F , and B .
- Screen size in pixels R and C .

As with parallel projection, the basic process is to align the view reference coordinate system with the world coordinate axes, move the COP to the origin and shear to center the view volume, scale the volume to an appropriate size, and then project the world onto a plane using a perspective projection.

The process:

1. Calculate the U vector for the view reference coordinate system

$$U = \text{VUP} \times \text{VPN} \quad (58)$$

2. Recalculate the VUP so that the view reference coordinates are orthogonal

$$\text{VUP} = \text{VPN} \times U \quad (59)$$

3. Translate the VRP to the origin

$$\text{VTM}_1 = T(-\text{VRP}) \quad (60)$$

4. Normalize U , VUP and VPN and use the R_{xyz} rotation method to align the axes.

$$\text{VTM}_2 = R_{xyz}(\hat{U}, \hat{\text{VUP}}, \hat{\text{VPN}})\text{VTM}_1 \quad (61)$$

5. Translate the COP to the origin. In many cases, the COP is just defined by the projection distance d . In that case, the translation becomes $T(0, 0, d)$.

$$\text{VTM}_3 = T(-\text{COP}_x, -\text{COP}_y, -\text{COP}_z)\text{VTM}_2 \quad (62)$$

6. Calculate the direction of projection (in case the COP or window center are not along the VPN).

$$\text{DOP} = \begin{bmatrix} \frac{u_{\min} + u_{\max}}{2} - \text{COP}_x \\ \frac{v_{\min} + v_{\max}}{2} - \text{COP}_y \\ -\text{COP}_z \end{bmatrix} \quad (63)$$

7. Shear by the DOP. Note that if the COP is located along the $-\text{VPN}$, then no shearing is required.

$$\text{VTM}_4 = Sh_z\left(\frac{D_x}{D_z}, \frac{D_y}{D_z}\right)\text{VTM}_3 \quad (64)$$

8. Scale to the canonical view volume, which is a pyramid.

- Update the location of the VRP

$$\text{VRP}' = \text{VTM}_4 \text{VRP} \quad (65)$$

- Update the distance from the origin to the back clip plane.

$$B' = \text{VRP}'_z + B \quad (66)$$

- The width of the view window is $du = u_1 - u_0$
- The height of the view window is $dv = v_1 - v_0$

$$\text{VTM}_{\text{World} \rightarrow \text{CVV}} = S\left(\frac{2\text{VRP}'_z}{B'du}, \frac{2\text{VRP}'_z}{B'dv}, \frac{1}{B'}\right) \text{VTM}_4 \quad (67)$$

The canonical view volume for perspective projection is a pyramid with its apex at the origin. The back clip plane is at $Z = 1$, and the base of the pyramid is a 2×2 square from $(-1, -1)$ to $(1, 1)$ in X and Y . As with parallel projection, clipping needs to occur here, particularly front and back clipping. The front clip plane is located at F' .

$$F' = \frac{\text{VRP}'_z + F}{B'} \quad (68)$$




Figure 7: The scale parameters necessary to convert the view volume to the canonical pyramid. When the COP is not on the negative VPn, then VRP'_z is used as the distance d .

9. Project the scene onto the view plane, located a distance d along the VPn.

$$d = \frac{\text{VRP}'_z}{B'} \quad (69)$$

$$P_{\text{per}} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & \frac{1}{d} & 0 \end{bmatrix} \quad (70)$$

10. Scale to the image size, negate the x - and y -axis, and translate the origin to the center of the image. Note that the view plane is $2d \times 2d$.

$$\text{VTM}_{\text{CVV} \rightarrow \text{Screen}} = T\left(\frac{C}{2}, \frac{R}{2}\right) S\left(-\frac{C}{2d}, -\frac{R}{2d}\right) P_{\text{per}} \quad (71)$$

5.5.1 Perspective Projection Matrix

Why does the perspective projection matrix work? Consider figure 8. The projection of the point y onto the image plane produces $y' = \frac{yd}{z}$.




Figure 8: Similar triangles show how the scene location is converted to a point on the view plane.

Now consider the effect of the perspective matrix on the point $[xyz1]'$. Re-normalizing the homogeneous point after the transformation produces the perspective projection onto the view plane.

$$\begin{bmatrix} \frac{xd}{z} \\ \frac{yd}{z} \\ \frac{d}{z} \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ y \\ z \\ \frac{d}{z} \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & \frac{1}{d} & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} \quad (72)$$

Therefore, the perspective transformation requires us to re-normalize the homogeneous coordinates after multiplying by the perspective projection matrix. One caveat to the process, however, is that it can be important to leave the z -coordinate alone rather than also dividing it by h . The perspective transformation converts all z -coordinates to the value d , but we may still have a use for the z -coordinate later on when determining visibility. In practice, the following pipeline is useful to implement.

1. Multiply by $VTM_{\text{World} \rightarrow \text{CVV}}$ to obtain CVV coordinates.
2. Clip graphics objects to the front and back clip planes.
3. Multiply by $VTM_{\text{CVV} \rightarrow \text{Screen}}$.
4. Divide x and y by the homogeneous coordinate to obtain the true screen coordinates.
5. Draw the graphics object to the screen.

5.5.2 Simple Perspective View

There are two assumptions we can make that simplify the view transformation process.

- The VRP is the center of the view window, and the window is defined by its width du and height dv .
- The COP is on the negative VPN and is defined by a distance d .

With the simplifying assumptions, the new process is as follows.

1. Calculate the U vector for the view reference coordinate system
2. Recalculate the VUP so that the view reference coordinates are orthogonal
3. Translate the VRP to the origin
4. Normalize U , VUP and VPN and use the R_{xyz} rotation method to align the axes.
5. Translate the COP to the origin using the transformation $T(0, 0, d)$.

Update the back clip plane distance $B' = d + B$

6. Scale to the canonical view volume using $S(\frac{2d}{B'd_u}, \frac{2d}{B'd_v}, \frac{1}{B'})$
7. Project the scene onto the view plane, located a distance d' along the VPN.

$$d' = \frac{d}{B'} \quad (73)$$

8. Scale to the image size, negate the x- and y-axis, and translate the origin to the center of the image.

Figure 9 shows all of the parameters of the simple perspective view in world coordinate space.




Figure 9: The complete specification of the perspective view.

5.6 Viewing Pipeline

Now that we know the steps required to move from 3D models to a 2D image, we can build a rough model of the viewing pipeline. Modern graphics cards implement most of the steps in the pipeline, and do so using massive parallelism since the same process is applied to many, many data points. The pipeline assumes objects are in world coordinates and that the viewing geometry is defined.

1. Transform the graphics primitives—usually points, lines, or triangles—to the canonical view volume
2. Clip the graphics primitives to the CVV
3. Transform the graphics primitives into screen space
4. Draw the visible primitives into the image

We have developed techniques for doing everything except the last step. Hidden surface removal (which means the same thing as visible surface drawing) is a topic we'll cover in section 8.

There are some basic CS concepts tucked into the graphics pipeline. The first is whether the pipeline should be applied serial to each graphics primitive or whether all of the primitives should move along the pipeline simultaneously. For transformations and clipping, it turns out there isn't much difference between the two approaches. Moving the data along each step simultaneously offers some potential caching advantages, but the number of operations is similar.

When drawing the visible primitives onto the screen, however, there are definite benefits to handling all of the graphics primitives simultaneously. The primary benefit is that we can avoid drawing over a pixel multiple times. Instead, if we have access to all of the polygons that cover a particular pixel, we can identify which polygon should be drawn and color the pixel once. It also permits us to more effectively implement concepts like anti-aliasing and transparency. We'll go over these algorithms and concepts when we cover hidden surface removal.

6 Hierarchical Modeling

A scene is a complex combination of basic primitive shapes. Most real-time rendered graphics, for example, consist purely of triangles. However, building a complex scene from triangles is similar to building a complex image one pixel at a time. Scenes are not, in general, abstract combinations of shapes, but are constructed from objects that have coherence. When we want a table in a scene, we do not think of it as 1000 triangles, but as a single object. We want to be able to place, orient, and scale the table as a unit.

Hierarchical modeling allows us to construct a scene out of parts. Each part may be further subdivided into simpler parts: a table might be constructed from four legs and a top. Having a hierarchical model of the scene permits us to separately manipulate parts of the scene at different levels. We can move the table as a unit, or we could expand the size of the legs, or we could add a small mark to one polygon.

At the root of hierarchical modeling is the concept of a structure or module. A module is one or more primitive graphic elements, attributes, and local transformations, organized as a directed acyclic graph. In addition to graphics primitives and transformations, a module can include other modules, enabling a module to become a node in a larger graph.

- The graph defines a sequence of operations that generate drawing commands
- Modules can contain modules
- Traversal of the graph is depth first
- A module can contain another module many different times (multiple instances)

It's important to understand that a module is a recipe for drawing something. When a module is inserted into a scene graph then the scene contains an **instance** of the module. A scene can contain many instances of a module, all of which have identical recipes.

One method of building the scene graph is, in fact, to have only a single physical copy of each module, storing a reference to the module in the overall scene graph for each instance. Making a copy of a module is more costly, but enables modifications to a single instance that do not propagate to the rest of the scene. However, one of the strengths of using a reference instead of a copy is that changes to a module do propagate across a scene with no further work required. If you find a missing polygon in your chair model, for example, changing your code in one place fixes it for the entire scene.

Modules allow us to build up a scene in layers that correspond to logical partitions. They reduce duplication and enable fast development of complex scenes with a small number of actual 3D polygons. A unit box module, for example, can be used to build an arbitrary box of any dimension by preceding the box with a set of transformations. More generally, the combination of a set of transformations and simple geometric primitives—e.g. spheres, cylinders, boxes, or pyramids—provides the basic building blocks for many objects and scenes.

Within a module, the elements form an ordered list. When traversing the list, the interpreter must maintain a current transformation state. Any graphics primitive is drawn using the current transformation state, generally called the **local transformation matrix** or **[LTM]**. The local transformation matrix collects transformations by pre-multiplying the most recent transformation by itself.

$$LTM_i = T_i LTM_{i-1} \quad (74)$$

A module must also take into account the **global transformation matrix** or [GTM] from its parent module. When the interpreter finds a module in the list of elements, it right multiplies the current LTM with the GTM and passes the product to the child module as its GTM.

$$GTM_{child} = GTM_{parent} LTM_{parent} \quad (75)$$


Example: A Robot Arm

A robot arm (or human arm) is an example of an object with a hierarchical relationship. If the elbow joint moves, then the hand moves as well. If the shoulder moves, the whole arm moves.


We can build up a robot arm using an end effector (a triangle) and a lines.

1. Create the hand module
 - Draw line 1
 - Draw line 2
 - Draw line 3
2. Create the forearm
 - Draw a line for the forearm
 - Rotate by the end-effector angle $R_z(\theta_1)$
 - Translate to the end of the forearm $T(d_f, 0)$
 - Draw the hand module
3. Create the upper arm
 - Draw a line for the upper arm
 - Rotate by the forearm angle $R_z(\theta_2)$
 - Translate to the end of the upper arm $T(d_u, 0)$
 - Draw the forearm module
4. Create the base
 - Rotate by the upper arm angle $R_z(\theta_3)$
 - Draw the upper arm

Each of the joints is encoded as a rotation parameter. To figure out what the transformation is for any given line, we simply concatenate the transformations from the node to the root. By adjusting the rotation angles, we can make the robot move in what appears to be a realistic manner. The entire arm is parameterized by the joint angles and the length of the translations.



(a)



(b)

Figure 10: (a) Robot arm with three joints and (b) its associated scene graph.

6.1 Building a Scene Graph

How do we implement the mechanics of building a scene graph? What capabilities do we need?

- We need to be able to begin and end modules.
- We need to be able to add graphics primitives to modules (points, lines, polygons, circles, ellipses, polylines, etc)
- We need to be able to add transformations (scales, rotations, translations)
- We could also use the module structure to manage drawing attributes (OpenGL does this)

In a very real sense we are designing a programming language. The keywords for our programming language could be:

- begin - opens a new module for collecting commands
- end - terminates collecting commands to the currently open module
- translate - adds a translation to the open module
- rotate - adds a rotation to the open module
- scale - adds a scale to the open module
- identity - sets the transformation to identity in the open module
- point - adds a point to the open module
- line - adds a line to the open module
- polygon - adds a polygon to the open module
- polyline - adds a polyline to the open module
- circle - adds a circle to the open module
- ellipse - adds an ellipse to the open module

A simple way to organize a module is using a linked list. In general, we will want to add things to the list quickly—so keep a link to the end of the list—and traverse the list. We will rarely do arbitrary insertion or deletions. Operations are always inserted at the end of the list.

There are many design decisions associated with building modules.

- Does the programmer get access to the modules?
 - OpenGL - programmer gets an ID number but no access to the underlying module list
 - Giving programmer access/control of modules permits more flexible programming
- How do you communicate the open module information to the routines that add primitives or transformations?
 - Could use a global variable that holds a pointer to the currently open module (or NULL if none exists)
 - Could pass in the module to which you want to add the item

- How do you represent the elements of a module (primitives, transformations, and modules)?
 - Need to encode what the primitive or transformation is
 - Need to encode the parameters, often different numbers of parameters
 - All need to be part of a linked list
 - Inheritance (C++) or union structures (C) are a reasonable way to represent elements.

A **union** in C is a polymorphic data structure that can hold one of a set of types. The compiler allocates enough space to hold the largest type, but the programmer controls which type is used to access it.

Example: Unions

We can use an Element as the node structures for our linked list, with a union to hold the object data.

```
// defines a union that can hold one instance of any of the constituent types
typedef union {
    Point point;
    Line line;
    Polygon polygon;
    Matrix matrix;
    void *module;
} Object;

// define an Element. The type field specifies what type the union holds
typedef struct {
    ObjectType type;
    Object obj;
    void *next;
} Element;

// define a Module as holding a linked list
typedef struct {
    Element *head;
    Element *tail;
} Module;
```

To access the fields of a union, use the field name in the union as a prefix to the structure's fields. Note that each union variable can hold only one of the given data types; it does not hold the different elements in different memory locations.

```
Object obj; // declare a union object
```

```
// treat the union as a point and access the x coordinate
obj.point.val[0] = 0.0;
```

```
// treat the union as a polygon and access the numVertex field
obj.polygon.numVertex = 4;
```

In the above code snippet, for example, setting numVertex means that trying to access the union as a point will no longer make sense.

6.2 Traversing a Scene Graph

We traverse a scene graph in order from the top level down in a depth-first fashion. During the traversal, we have to keep track of the current transformation matrix. It is useful to divide the transformations into three parts.

- VTM - View Transformation Matrix, affects all drawn elements
- GTM - Global Transformation Matrix, contains all of the transformations from the higher levels of the scene graph
- LTM - Local Transformation Matrix, contains all of the transformations seen so far in the current module

When traversing the graph, to draw a graphic primitive we always multiply the points by the current transformation matrix before drawing. The CTM is the matrix product of the VTM, GTM, and LTM.

$$[CTM] = [VTM][GTM][LTM] \quad (76)$$

When we encounter a transformation matrix, it always premultiplies the LTM. Conceptually, you can think of it as squeezing in between the GTM and LTM, which is why we need to keep them separate.

When we encounter a module, we need to create a new transformation matrix for the module that is the product of the GTM and LTM. Then we recursively call the scene traversal, passing in the new matrix as its GTM. When the module completes, the old GTM and LTM get restored. In most cases it makes sense to keep the VTM as a separate matrix throughout the graph traversal, passing it around as necessary.

The basic idea of the scene traversal algorithm is to loop over the elements of the current module's list and take the appropriate action depending upon the element's type.

- If the element is a graphics primitive, the algorithm calculates the current transformation and uses it to transform all of the points into the correct location. Then it draws the primitive using the transformed points.
- If the element is a matrix, the matrix pre-multiplies the LTM.
- If the element is a module, the algorithm calculates a new GTM, which is the old GTM multiplied by the LTM, and passes the module and the new GTM to the recursive call.

One potential addition to the list of operations above are **push matrix** and **pop matrix** operations for the local transformation matrix. OpenGL, for example, implements push and pop operations. This can be convenient when working with a module that has many pieces to it, all relative to the same location. The push operation sticks a copy of the current local transformation matrix (LTM) onto a stack. The pop operation restores the LTM to the top transformation matrix off the stack.

Scene Traversal Algorithm

The following is pseudo-code for the `Module.draw` function that shows the cases for a subset of the different graphics types.

```

Module_draw
  Parameters:
    Module md
    Matrix VTM
    Matrix GTM
    DrawState DS
    Lighting light
    Image src

  set the matrix LTM to identity

  for each element E in the module md
    switch on the type of E
      Color
        DS->color = color data in E
      Point
        copy the line data in E to X
        transform X by the LTM
        transform X by the GTM
        transform X by the VTM
        normalize X by the homogeneous coord
        draw X using DS->color (if X is in the image)
      Line
        copy the line data in E to L
        transform L by the LTM
        transform L by the GTM
        transform L by the VTM
        normalize L by the homogeneous coord
        draw L using DS->color
      Polygon
        copy the polygon data in E to P
        transform P by the LTM
        transform P by the GTM
        transform P by the VTM
        normalize P by the homogeneous coord
        if DS->shade is ShadeFrame
          draw the boundary of P using DS->color
        else if DS->shade is ShadeConstant
          draw P using DS->color
      Matrix
        LTM = (Matrix field of E) * LTM
      Identity
        LTM = I
      Module
        TM = GTM * LTM
        tempDS = DS
        Module_draw( (Module field of E), VTM, TM, tempDS, Light, src )

```